

UVA CS 4774: Machine Learning

S4: Lecture 21 Extra: (SVM Optimization and Dual basic)

Dr. Yanjun Qi

University of Virginia

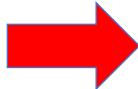
Department of Computer Science

Module IV
Extra

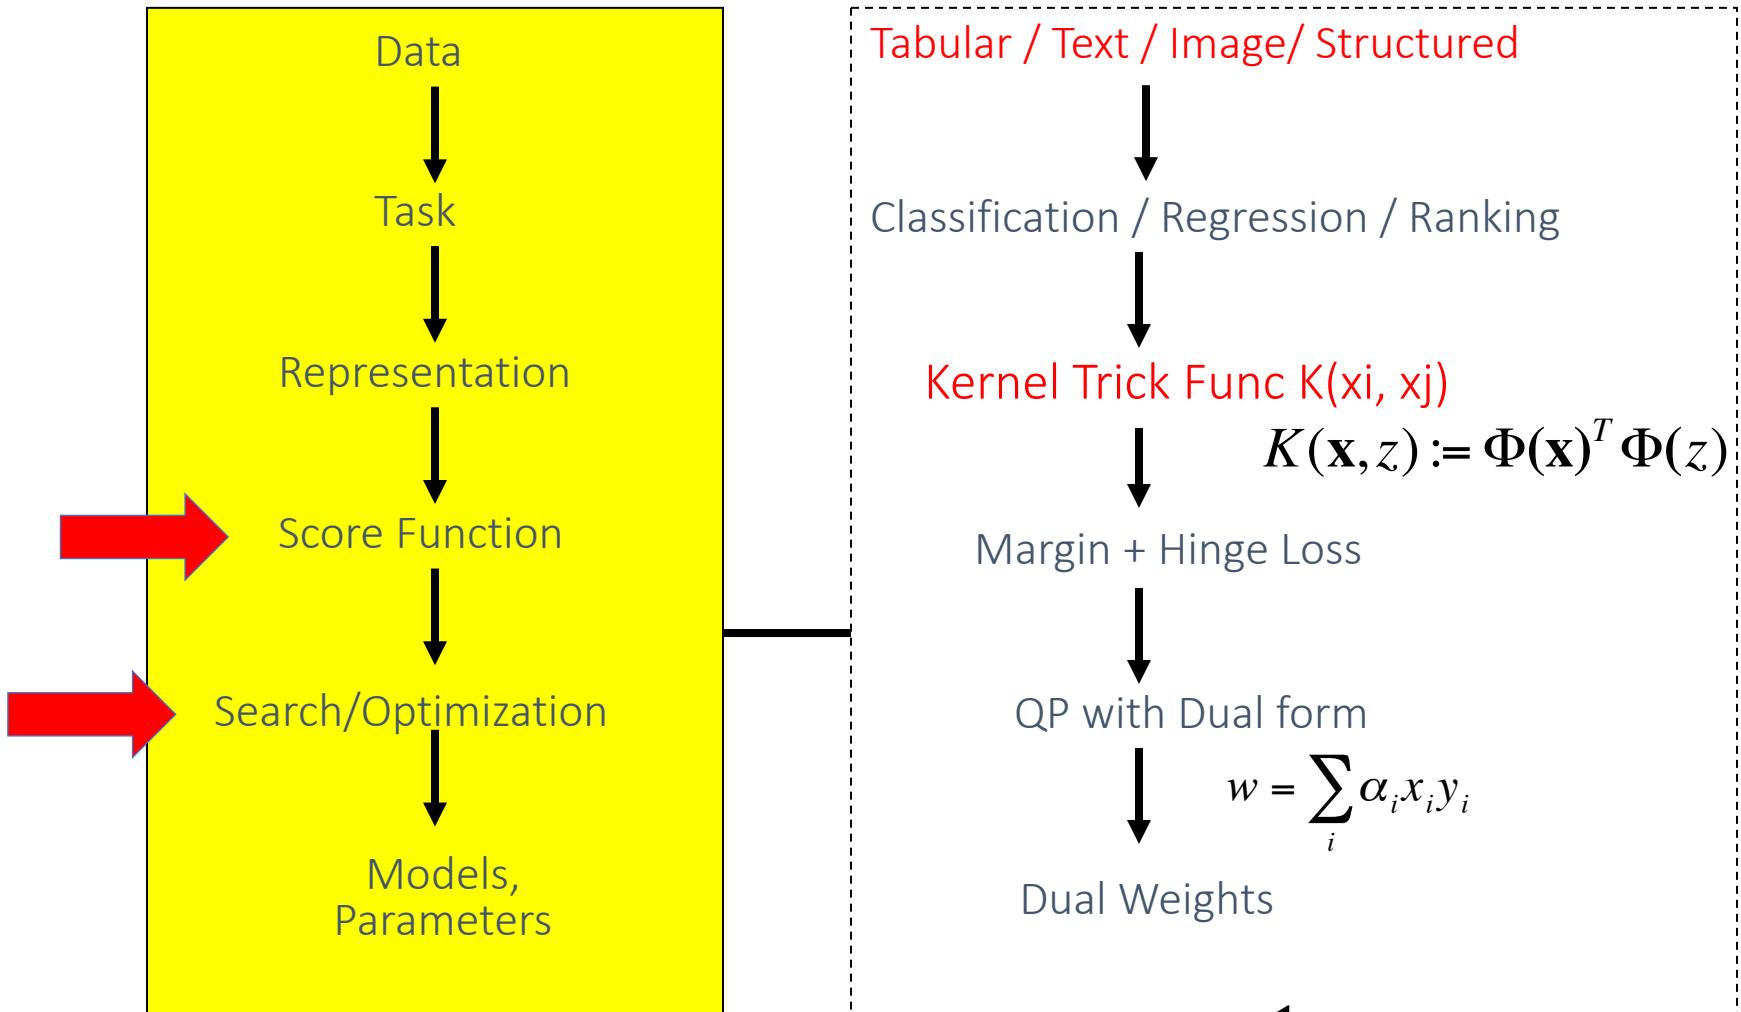
What Left in SVM?

- ❑ Support Vector Machine (SVM)
 - ✓ History of SVM
 - ✓ Large Margin Linear Classifier
 - ✓ Define Margin (M) in terms of model parameter
 - ✓ Optimization to learn model parameters (w, b)
 - ✓ Linearly Non-separable case (soft SVM)
 - ✓ Optimization with dual form
 - ✓ Nonlinear decision boundary
 - ✓ Practical Guide

$k(x, z)$



This: Kernel Support Vector Machine



$$\underset{\mathbf{w}, b}{\operatorname{argmin}} \sum_{i=1}^p w_i^2 + C \sum_{j=1}^n \varepsilon_j$$

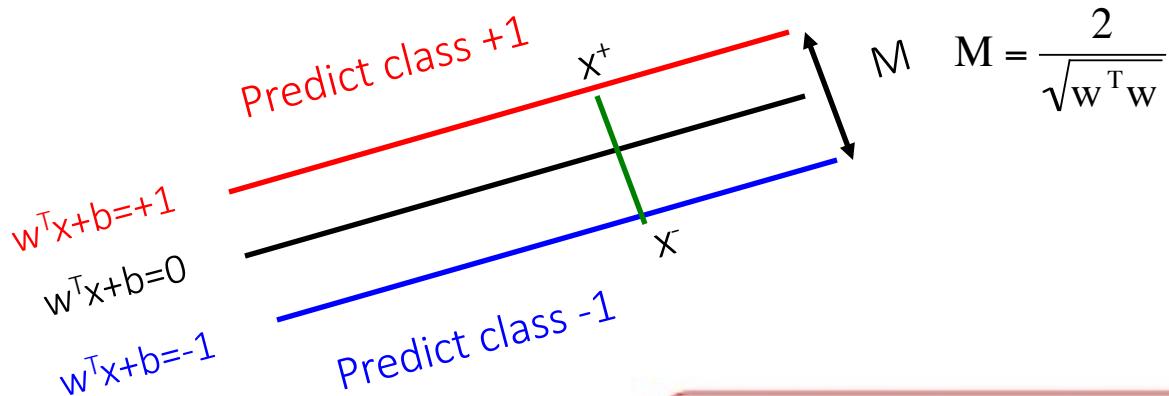
subject to $\forall \mathbf{x}_i \in Dtrain : y_i(\mathbf{x}_i \cdot \mathbf{w} + b) \geq 1 - \varepsilon_i$

$$\max_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j x_i^T x_j$$

$$\sum_i \alpha_i y_i = 0, \quad \alpha_i \geq 0 \quad \text{3} \quad \forall i$$

Optimization Step

i.e. learning optimal parameter for SVM



1. Correctly classifies all points
2. Maximizes the margin (or equivalently minimizes $w^T w$)

Optimization Reformulation

1. Correctly classifies all points
2. Maximizes the margin (or equivalently minimizes $w^T w$)

$$\text{Min } (w^T w)/2$$

subject to the following constraints:

For all x in class +1

$$w^T x + b \geq 1$$

$$y_i = 1$$

For all x in class -1

$$w^T x + b \leq -1$$

$$y_i = -1$$

A total of n
constraints if
we have n
input samples

$$\rightarrow \text{Pos } y_i = 1, w^T x_i + b \geq 1$$

$$y_i (w^T x_i + b) \geq 1$$

$$\rightarrow \text{Neg } y_i = -1, w^T x_i + b \leq -1$$

$$y_i (w^T x_i + b) \geq 1$$

Optimization Reformulation

1. Correctly classifies all points
2. Maximizes the margin (or equivalently minimizes $\mathbf{w}^T \mathbf{w}$)

$$\text{Min } (\mathbf{w}^T \mathbf{w})/2$$

subject to the following constraints:

For all x in class +1

$$\mathbf{w}^T \mathbf{x} + b \geq 1$$

For all x in class -1

$$\mathbf{w}^T \mathbf{x} + b \leq -1$$

}

A total of n constraints if we have n input samples



$$\underset{\mathbf{w}, b}{\operatorname{argmin}} \sum_{i=1}^p w_i^2 / 2$$

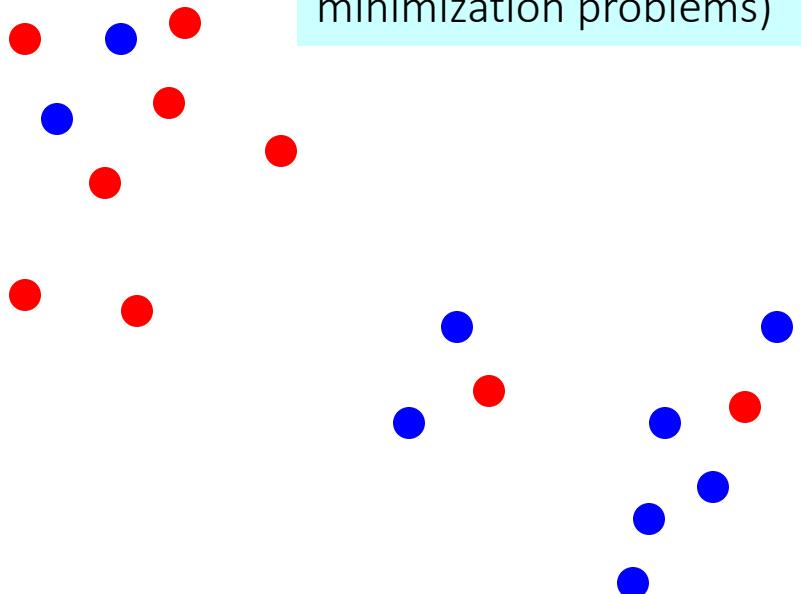
$$\text{subject to } \forall \mathbf{x}_i \in D_{\text{train}} : y_i (\mathbf{x}_i \cdot \mathbf{w} + b) \geq 1$$

Linearly Non separable case

- So far we assumed that a linear hyperplane can perfectly separate the points
- But this is not usually the case
 - noise, outliers

How can we convert this to a QP problem?

Hard to solve (two minimization problems)

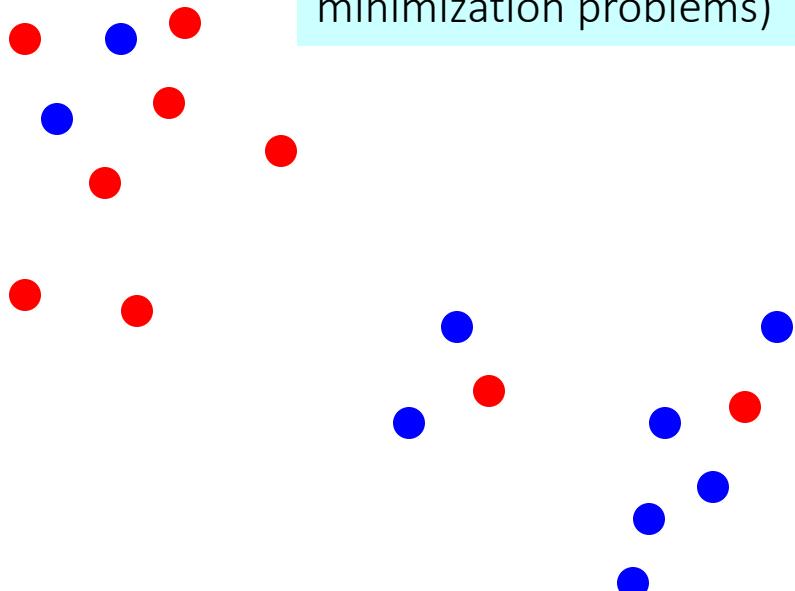


- Minimize training errors?

$$\left\{ \begin{array}{l} \min w^T w / 2 \\ \min \# \text{errors} \end{array} \right.$$

Linearly Non separable case

- So far we assumed that a linear plane can perfectly separate the points
- But this is not usually the case
 - noise, outliers



How can we convert this to a QP problem?

- Minimize training errors?

$$\min w^T w / 2$$

$$\min \# \text{errors}$$

- Penalize training errors:

$$\min w^T w / 2 + C * (\# \text{errors})$$

Hard to encode in a QP problem

$C \uparrow$ penalize errors more

SVM : $\min_{\mathbf{w}, b} \mathbf{w}^T \mathbf{w} + \underbrace{C(\# \text{ errors})}_{\text{penalize errors}}$

Ridge Regression :

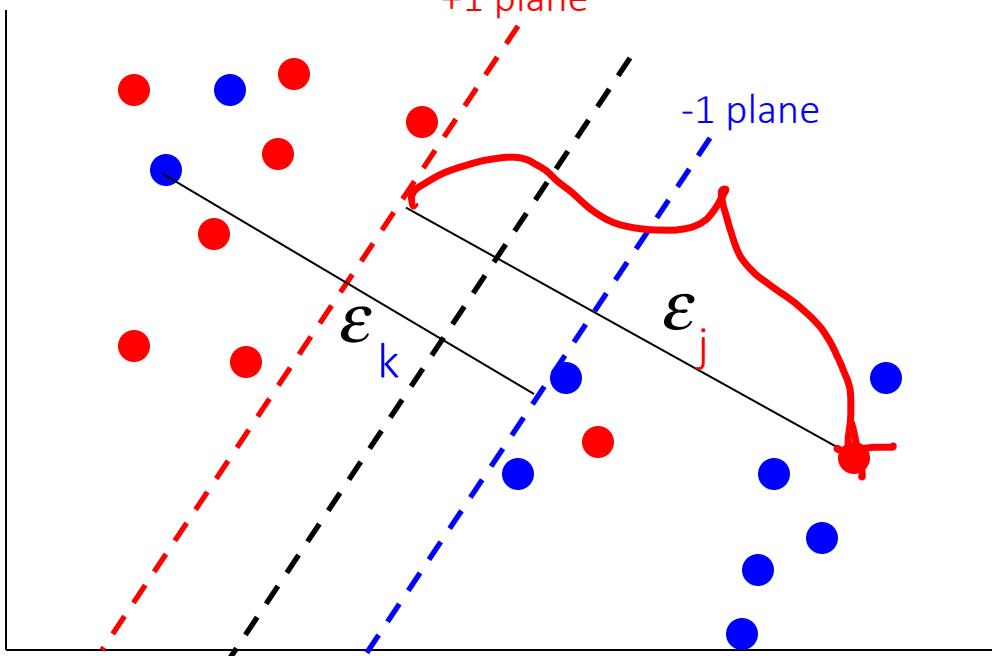
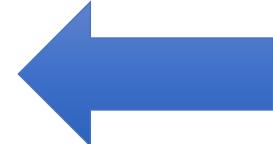
$$\min_{\theta} \lambda \theta^T \theta + J(\theta)$$

Linearly Non separable case

- Instead of minimizing the number of misclassified points we can minimize the distance between these points and their correct plane

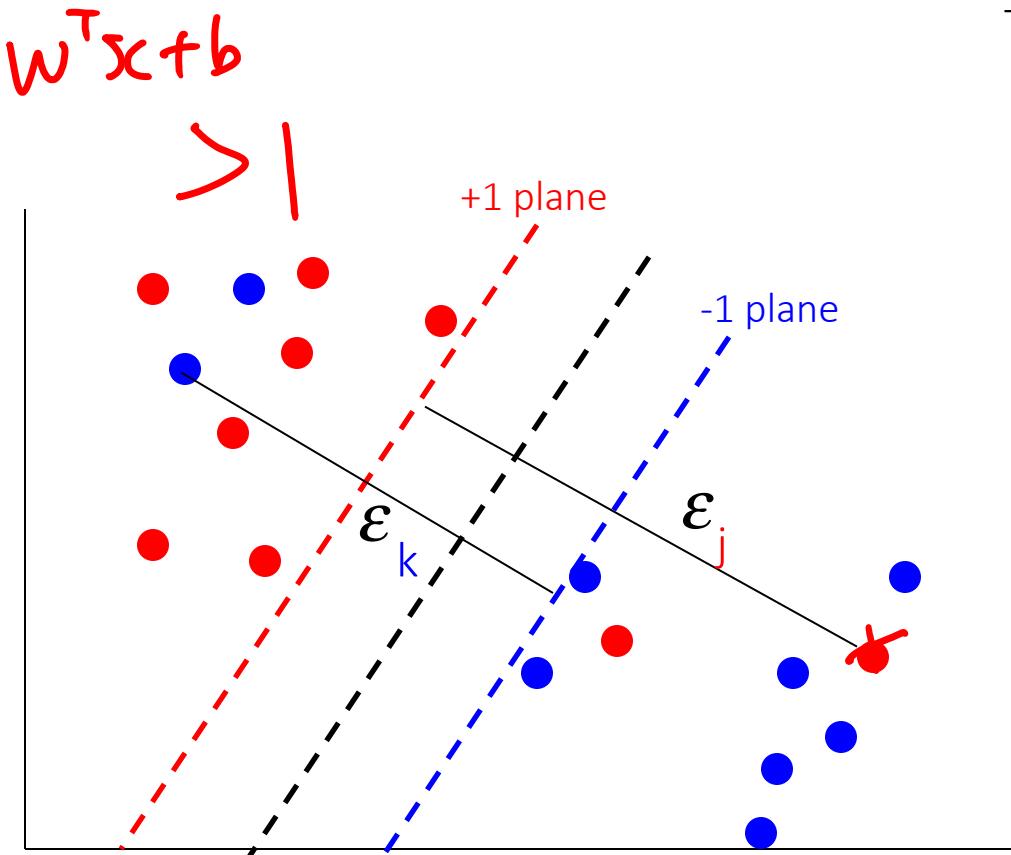
The new optimization problem is:

$$\min_w \frac{w^T w}{2} + C \sum_{i=1}^n \varepsilon_i$$



Linearly Non separable case

- Instead of minimizing the number of misclassified points we can minimize the distance between these points and their correct plane



The new optimization problem is:

$$\min_w \frac{w^T w}{2} + C \sum_{i=1}^n \varepsilon_i$$

subject to the following inequality constraint

For all x_i in class +1

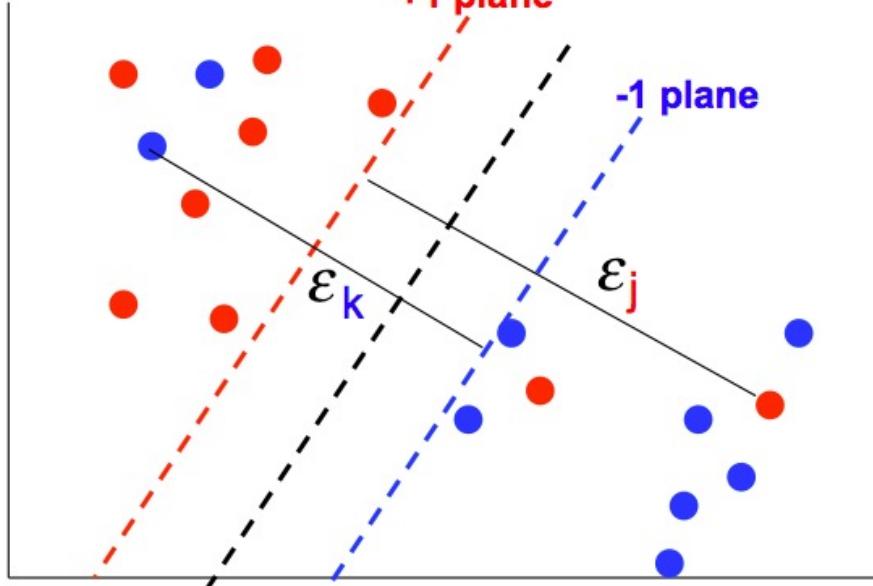
$$w^T x_i + b \geq 1 - \varepsilon_i \quad \varepsilon_i \geq 0$$

For all x_i in class -1

$$w^T x_i + b \leq -1 + \varepsilon_i$$

Wait. Are we missing something?

Final optimization for linearly non-separable case



The new optimization problem is:

$$\min_w \frac{w^T w}{2} + C \sum_{i=1}^n \varepsilon_i$$

subject to the following inequality constraints:

For all x_i in class +1

$$w^T x_i + b \geq 1 - \varepsilon_i$$

For all x_i in class -1

$$w^T x_i + b \leq -1 + \varepsilon_i$$

For all i

$$\varepsilon_i \geq 0$$

A total of n constraints

Another n constraints

Two optimization problems:

For the separable and non separable cases

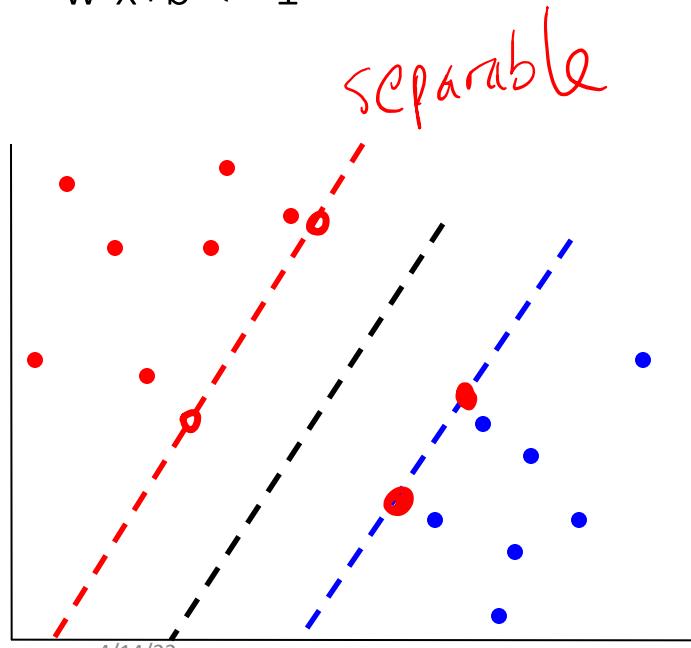
$$\min_w \frac{w^T w}{2}$$

For all x in class +1

$$w^T x + b \geq 1$$

For all x in class -1

$$w^T x + b \leq -1$$



$$\min_w \frac{w^T w}{2} + C \sum_{i=1}^n \varepsilon_i$$

For all x_i in class +1

$$w^T x_i + b \geq 1 - \varepsilon_i$$

For all x_i in class -1

$$w^T x_i + b \leq -1 + \varepsilon_i$$

For all i

Non Separable

$$\varepsilon_i \geq 0$$

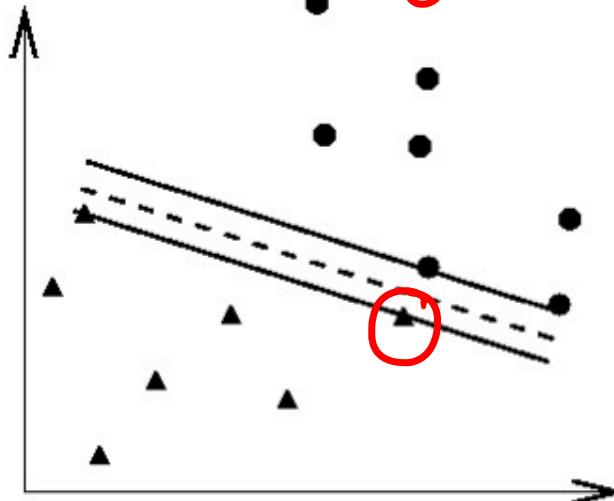
Model Selection, find right C

large C

Select the
right
penalty
parameter
 C

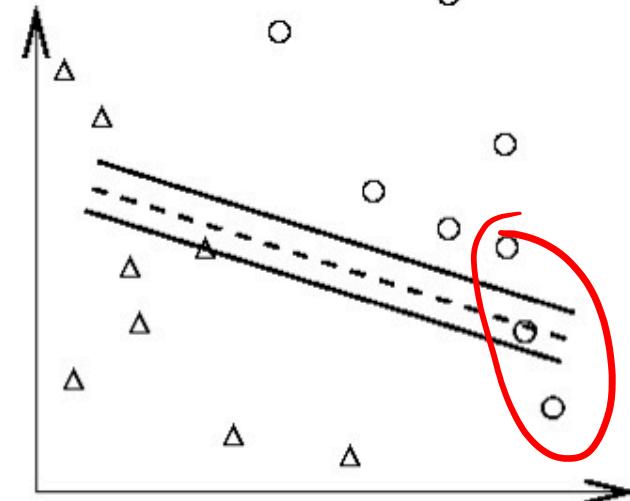
small C

Training

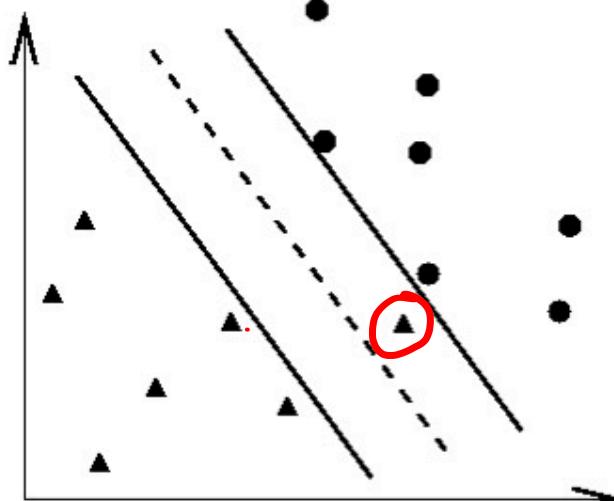


(a) Training data and an overfitting classifier

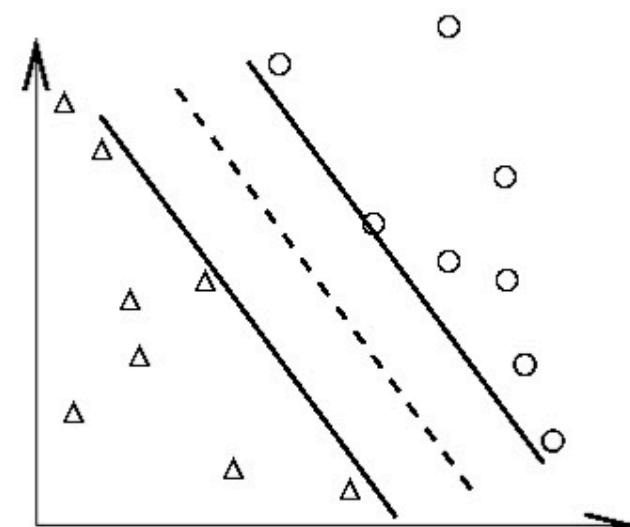
Test



(b) Applying an overfitting classifier on testing data



(c) Training data and a better classifier



(d) Applying a better classifier on testing data

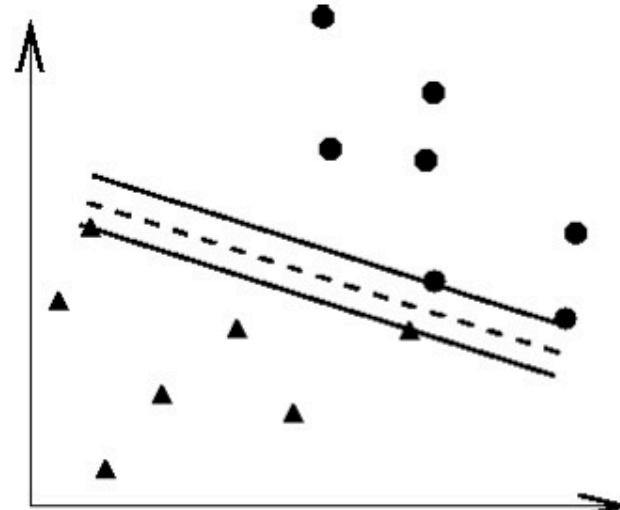
Model Selection, find right C

large C

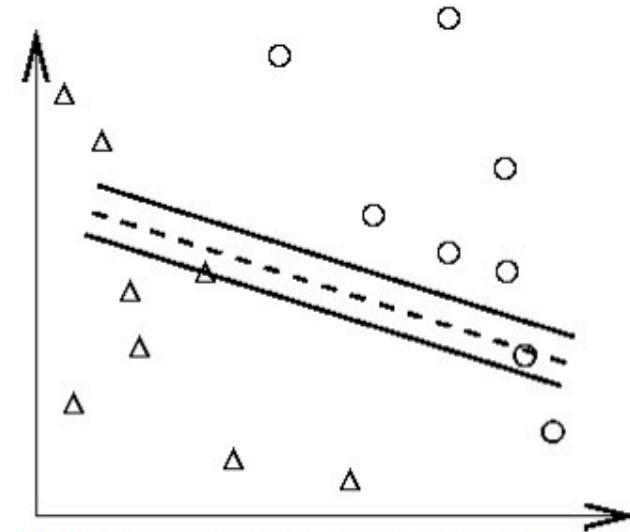
A large value of C means that misclassifications are bad - resulting in smaller margins and less training error (but more expected true error).

A small C results in more training error, hopefully better true error.

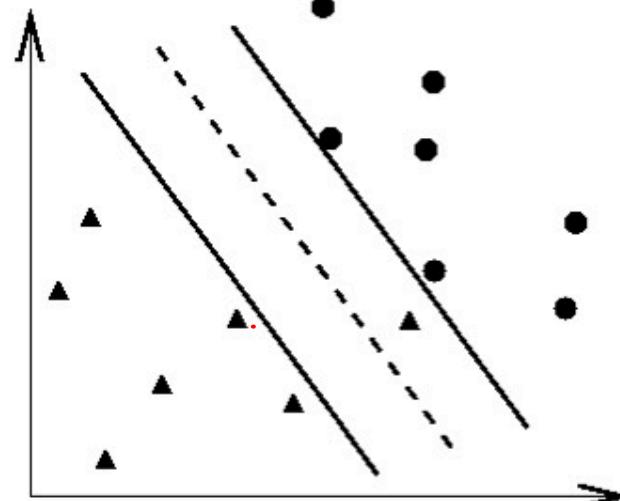
Small C



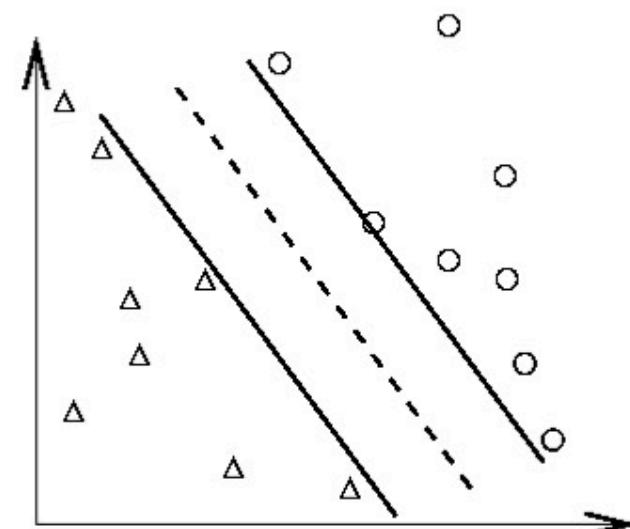
(a) Training data and an overfitting classifier



(b) Applying an overfitting classifier on testing data



(c) Training data and a better classifier



(d) Applying a better classifier on testing data

Hinge Loss for Soft SVM

$$\min_w \frac{\mathbf{w}^T \mathbf{w}}{2} + C \sum_{i=1}^n \varepsilon_i$$

For all x_i in class +1

$$\mathbf{w}^T \mathbf{x}_i + b \geq 1 - \varepsilon_i$$

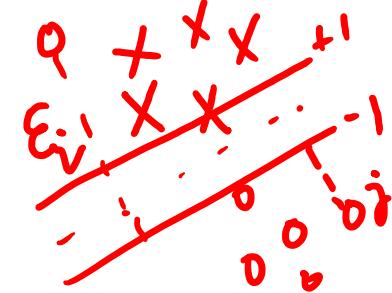
For all x_i in class -1

$$\mathbf{w}^T \mathbf{x}_i + b \leq -1 + \varepsilon_i$$

For all i

$$\varepsilon_i \geq 0$$

$$\sum_{i=1}^n \varepsilon_i$$



$$\sum_{i=1}^n \max(0, 1 - y_i f(\mathbf{x}_i))$$

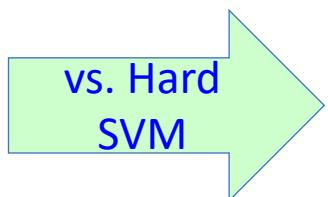


$$\operatorname{argmin}_{\mathbf{w}, b} \frac{\mathbf{w}^T \mathbf{w}}{2} + C \sum_{i=1}^n \max(0, 1 - y_i (\mathbf{w}^T \mathbf{x}_i + b))$$

subject to:

$$y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \varepsilon_i$$

$$\varepsilon_i \geq 0$$

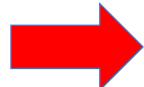


$$\operatorname{argmin}_{\mathbf{w}, b} \sum_{i=1}^p w_i^2 / 2$$

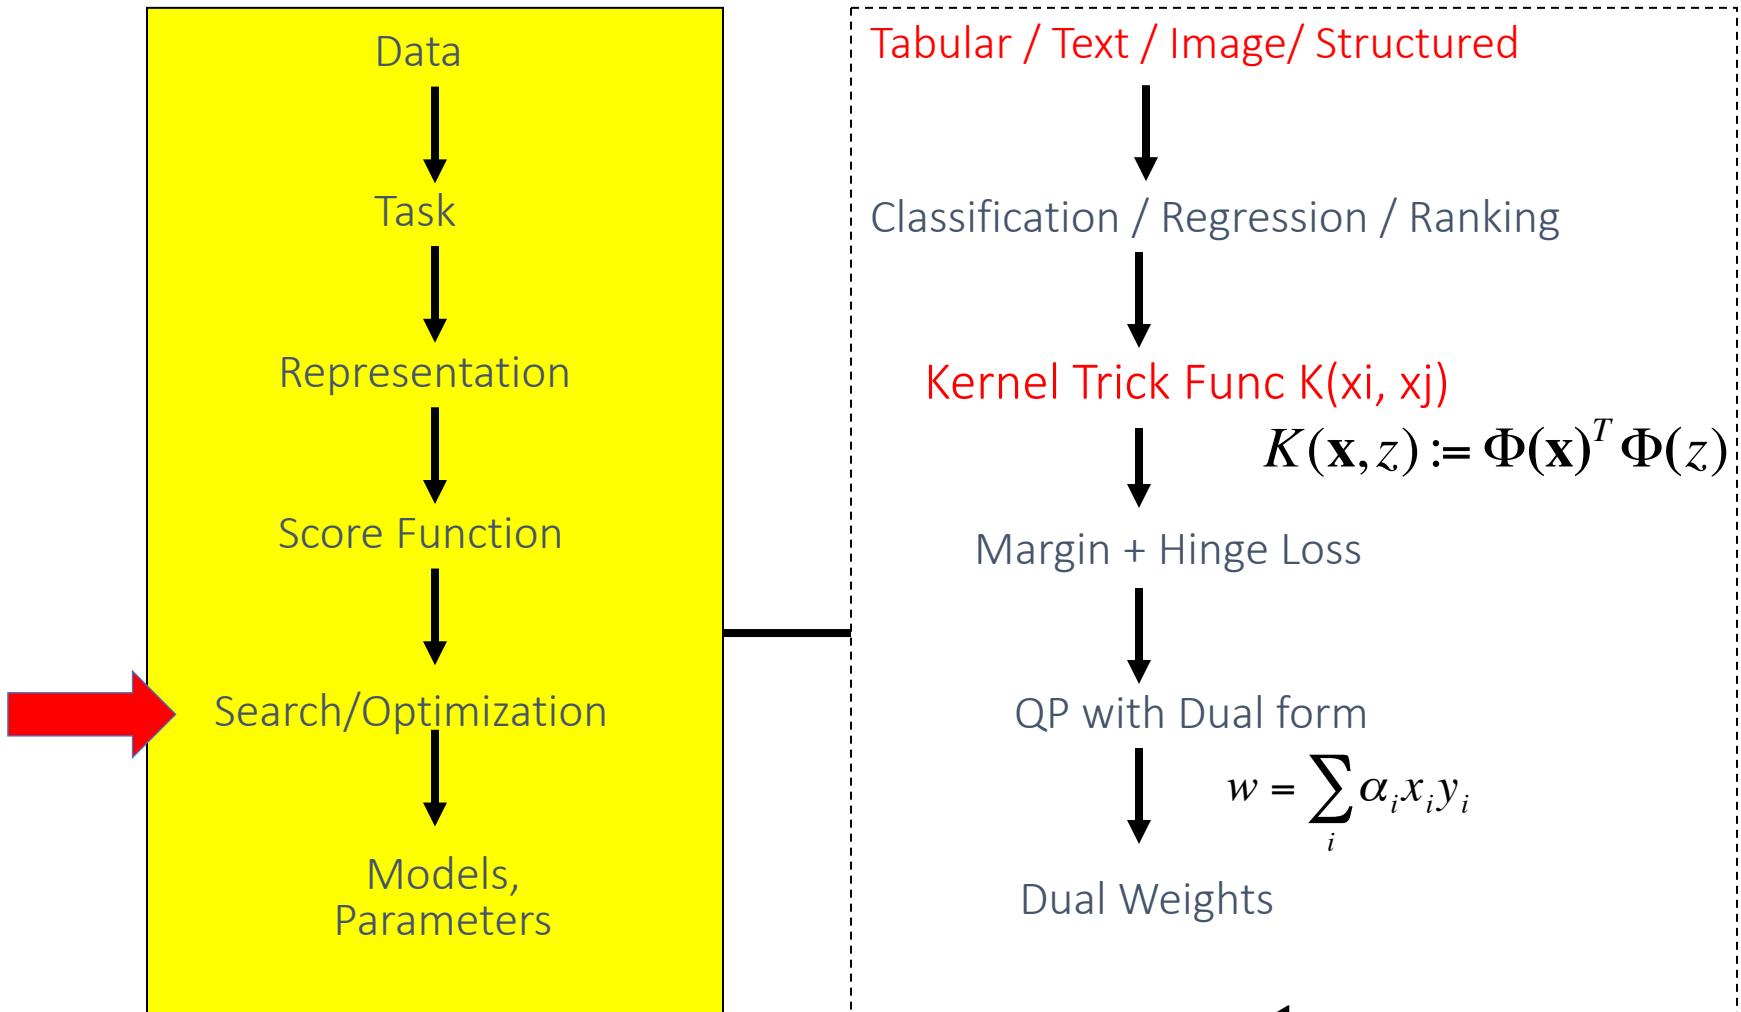
$$\text{subject to } \forall \mathbf{x}_i \in D_{train}: y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$$

What Left in SVM?

- ❑ Support Vector Machine (SVM)
 - ✓ History of SVM
 - ✓ Large Margin Linear Classifier
 - ✓ Define Margin (M) in terms of model parameter
 - ✓ Optimization to learn model parameters (w, b)
 - ✓ Linearly Non-separable case (soft SVM)
 - ✓ Optimization with dual form
 - ✓ Nonlinear decision boundary
 - ✓ Practical Guide



This: Kernel Support Vector Machine



$$\underset{\mathbf{w}, b}{\operatorname{argmin}} \sum_{i=1}^p w_i^2 + C \sum_{j=1}^n \varepsilon_j$$

subject to $\forall \mathbf{x}_i \in Dtrain : y_i(\mathbf{x}_i \cdot \mathbf{w} + b) \geq 1 - \varepsilon_i$

$$\max_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$$

$$\sum_j \alpha_j y_j = 0, \quad \alpha_j \geq 0 \quad \text{18} \quad \forall i$$

Two optimization problems: For the **separable** and **non separable** cases

$$\text{Min } (w^T w)/2$$

For all x in class + 1

$$w^T x + b \geq 1$$

For all x in class - 1

$$w^T x + b \leq -1$$

$$\min_w \frac{w^T w}{2} + C \sum_{i=1}^n \varepsilon_i$$

For all x_i in class + 1

$$w^T x_i + b \geq 1 - \varepsilon_i$$

For all x_i in class - 1

$$w^T x_i + b \leq -1 + \varepsilon_i$$

For all i

$$\varepsilon_i \geq 0$$

- Instead of solving these QPs directly we will solve **a dual formulation of the SVM optimization problem**
- The main reason for switching to this type of representation is that it would allow us to use a **neat trick that will make our lives easier** (and the run time faster)

Optimization Review: Ingredients

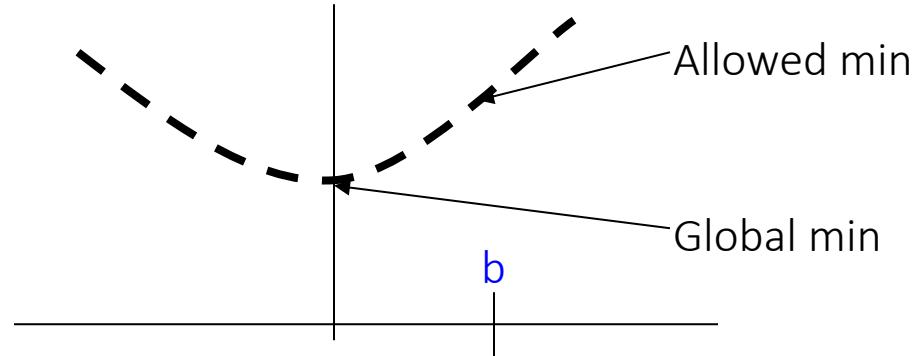
- Objective function
- Variables
- Constraints

**Find values of the variables
that minimize or maximize the objective function
while satisfying the constraints**

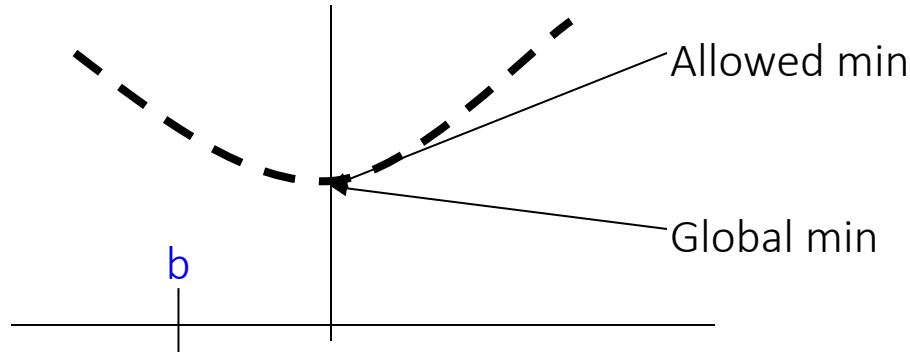
Optimization Review: Constrained Optimization

$$\begin{aligned} & \min_u u^2 \\ \text{s.t. } & u \geq b \end{aligned}$$

Case 1:



Case 2:

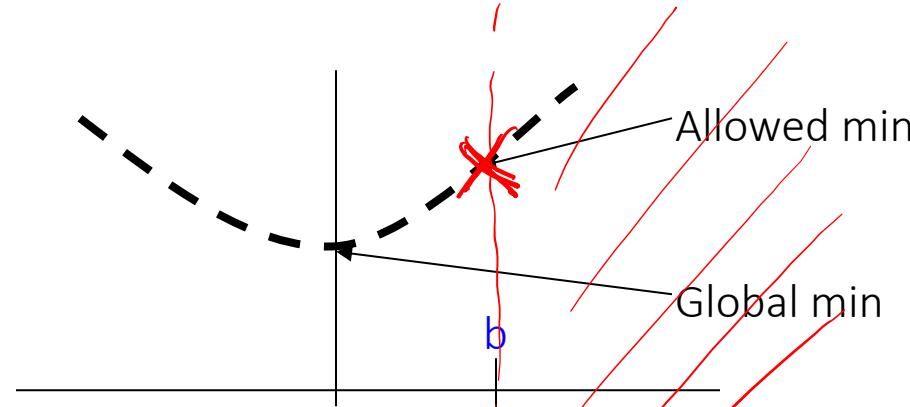


Optimization Review: Constrained Optimization

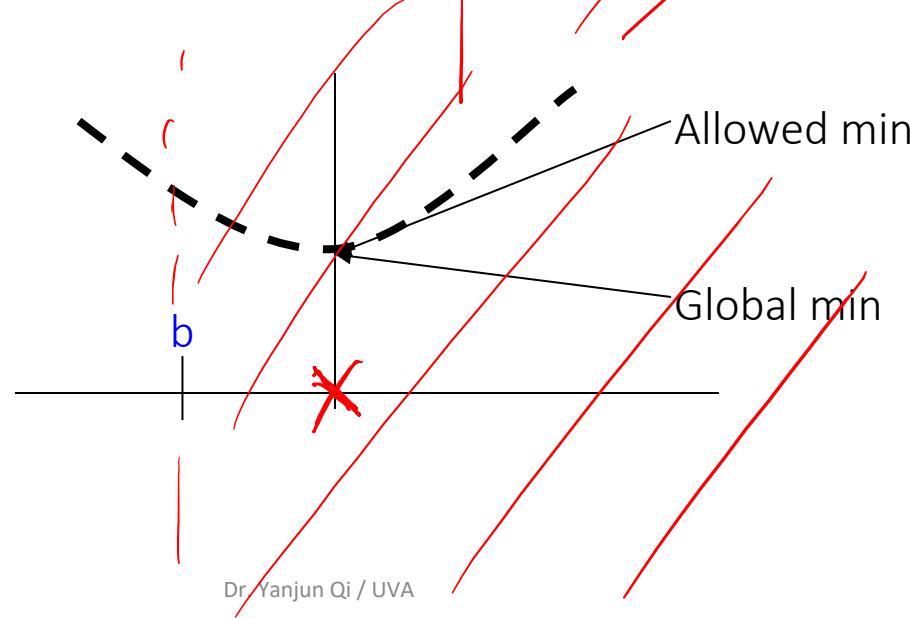
$$f(u) = u^2$$

$$\begin{aligned} & \min_u u^2 \\ & \text{s.t. } u \geq b \end{aligned}$$

Case 1:



Case 2:



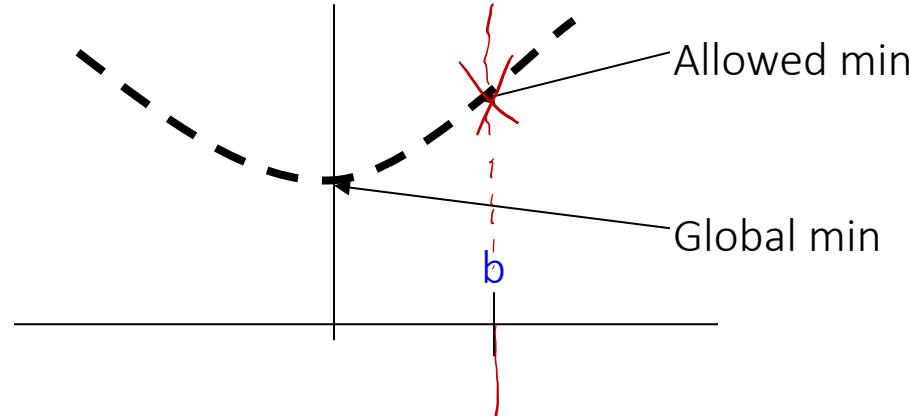
Optimization Review: Constrained Optimization

$$f(u)$$

$$\begin{aligned} & \min_u u^2 \\ & \text{s.t. } u \geq b \end{aligned}$$

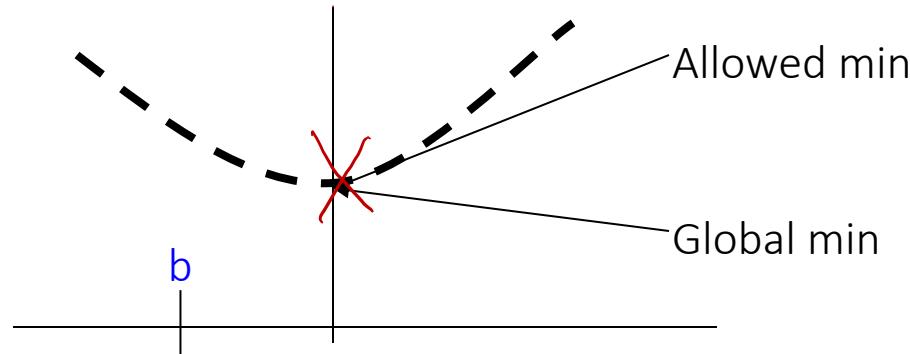
\curvearrowleft Subject to

Case 1:



$$\begin{cases} b > 0 \\ f(u^*) = b^2 \\ u^* = b \end{cases}$$

Case 2:



$$\begin{cases} b < 0 \\ f(u^*) = 0 \\ u^* = 0 \end{cases}$$

$$\begin{array}{l} \min_u u^2 \\ \text{s.t. } u \geq b \end{array}$$

$$\left\{ \begin{array}{l} \min_u f_0(u) = u^2 \\ \text{s.t. } b - u \leq 0 \end{array} \right.$$

Primal
problem

$$\begin{array}{l} \min_u u^2 \\ \text{s.t. } u \geq b \end{array}$$

$$\left\{ \begin{array}{l} \min_u f_0(u) = u^2 \\ \text{s.t. } b - u \leq 0 \end{array} \right.$$

\Rightarrow multiplier variable

$$② L(u, \alpha) = \underbrace{u^2}_{\geq 0} + \underbrace{\alpha(b-u)}_{\leq 0}$$

$$\begin{array}{l} \min_u u^2 \\ \text{s.t. } u \geq b \end{array}$$

$$\text{① } \left\{ \begin{array}{l} \min_u f_0(u) = u^2 \\ \text{s.t. } b - u \leq 0 \end{array} \right.$$

$$\text{② } L(u, \alpha) = \underbrace{u^2}_{\substack{\downarrow \\ (x)}} + \underbrace{\alpha(b-u)}_{\substack{\geq 0 \\ \leq 0}}$$

$$\text{③ } \frac{\partial L(u, \alpha)}{\partial u} = 2u - \alpha = 0$$

$$u = \frac{\alpha}{2}$$

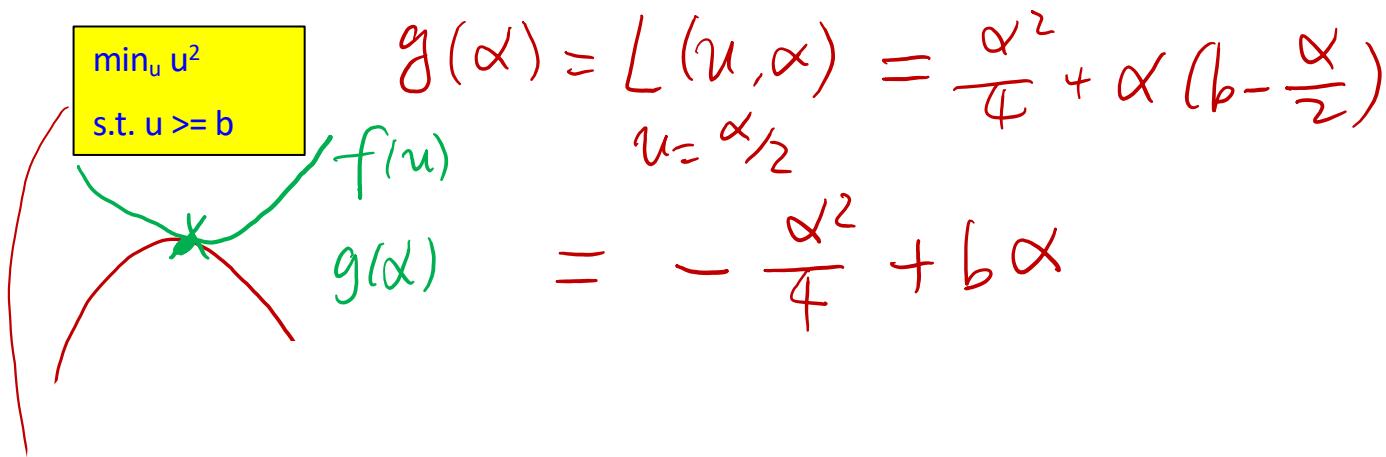
$$\rightarrow \arg \min_u L(u, \alpha)$$

$$\begin{array}{l} \min_u u^2 \\ \text{s.t. } u \geq b \end{array}$$

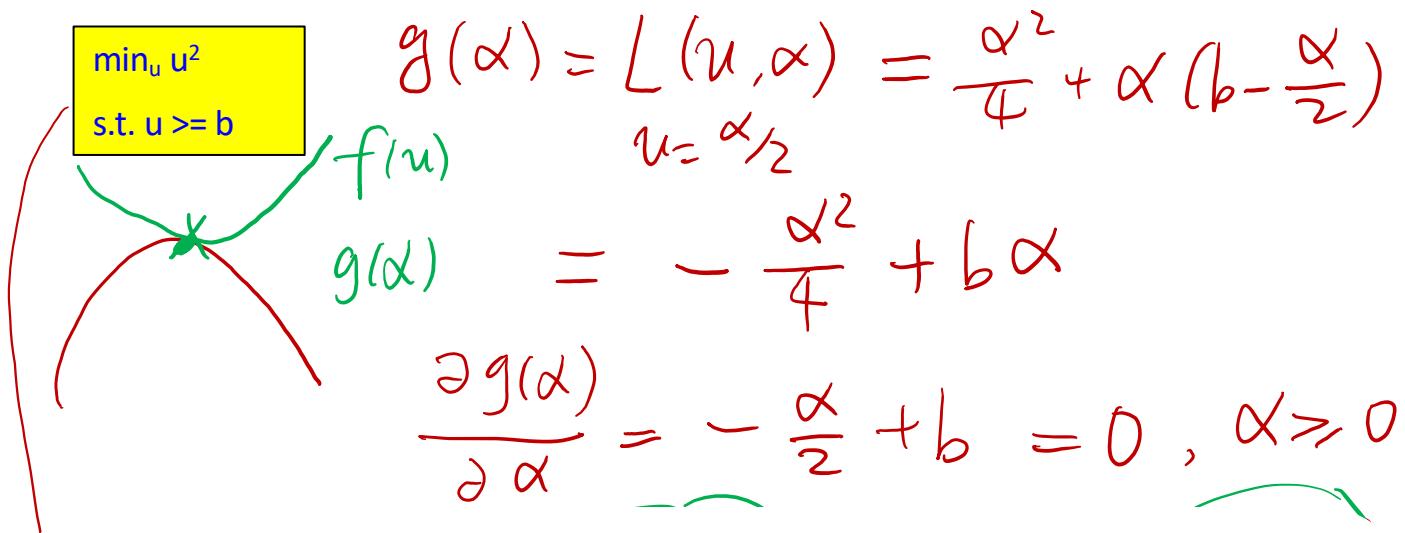
$$g(\alpha) = L(u, \alpha) = \frac{\alpha^2}{4} + \alpha(b - \frac{\alpha}{2})$$



$$\begin{array}{l} \min_u u^2 \\ \text{s.t. } u \geq b \end{array}$$



$$\begin{array}{l} \min_u u^2 \\ \text{s.t. } u \geq b \end{array}$$



$$\begin{array}{ll} \min_u u^2 \\ \text{s.t. } u \geq b \end{array}$$

$$g(\alpha) = L(u, \alpha) = \frac{\alpha^2}{4} + \alpha(b - \frac{\alpha}{2})$$

$f(u)$

$g(\alpha)$

$u = \alpha/2$

$$= -\frac{\alpha^2}{4} + b\alpha$$

$$\frac{\partial g(\alpha)}{\partial \alpha} = -\frac{\alpha}{2} + b = 0, \alpha \geq 0$$

\Rightarrow

Dual

$$\left\{ \begin{array}{l} b > 0 \\ b < 0 \end{array} \right.$$

$$\alpha = 2b, g(\alpha) = b^2$$

$$\alpha = 0, g(\alpha) = 0$$

$$\left\{ \begin{array}{l} b > 0 \\ b < 0 \end{array} \right. , g(\alpha) = b^2$$

\Rightarrow

Primal

$$\left\{ \begin{array}{l} b > 0 \\ b < 0 \end{array} \right. , f(u) = b^2$$

$$f(u) = 0$$

$$u = b$$

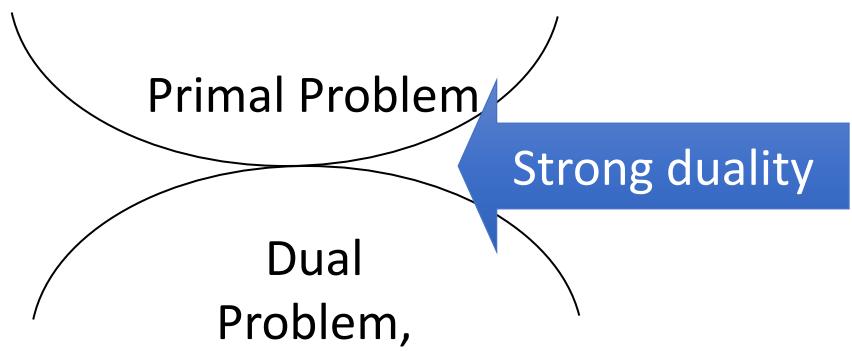
$$u = 0$$

$$\text{Primal} : \min_w \max_{\alpha} L(w, \alpha)$$

$$\text{Dual} : \underbrace{\max_{\alpha} \min_w L(w, \alpha)}_{.}$$
$$\Rightarrow \max_{\alpha} g(\alpha)$$

Optimization Review: Dual Problem (Extra)

- Solving dual problem if the dual form is easier than primal form
- Need to change primal **minimization** to dual **maximization** (OR → Need to change primal **maximization** to dual **minimization**)
- Only valid when the original optimization problem is convex/concave (strong duality)



Concrete derivation in L11Extra-SVMoptimalDual

$$f(u) : \begin{cases} \min u^2 \\ \text{s.t. } u \geq b \end{cases}$$

$$g(\alpha) : \begin{cases} \max -\frac{\alpha^2}{4} + b\alpha = \max \left\{ -\underbrace{(\alpha-b)^2}_{\geq 0} + b^2 \right\} \\ \text{s.t. } \alpha \geq 0 \end{cases}$$

$$\begin{cases} \text{if } b \geq 0, \quad u^* = b, \quad g^* = b^2 \\ \text{if } b < 0, \quad \alpha^* = 0, \quad g^* = 0 \end{cases}$$

$\Rightarrow \alpha(b-u) = 0$ KKT condition

Optimization Review: Lagrangian Duality (Extra)

- The Primal Problem

$$\min_w f_0(w)$$

Primal:

$$\text{s.t. } f_i(w) \leq 0, \quad i = 1, \dots, k$$

The generalized Lagrangian:

$$L(w, \alpha) = f_0(w) + \sum_{i=1}^k \alpha_i f_i(w)$$

the α 's ($\alpha_i \geq 0$) are called the Lagrangian multipliers

Lemma:

$$\max_{\alpha, \alpha_i \geq 0} L(w, \alpha) = \begin{cases} f_0(w) & \text{if } w \text{ satisfies primal constraints} \\ \infty & \text{o/w} \end{cases}$$

A re-written Primal:

$$\min_w \max_{\alpha, \alpha_i \geq 0} L(w, \alpha)$$

“Method of Lagrange multipliers”
convert to a higher-dimensional problem

Optimization Review: Lagrangian Duality, cont. (Extra)

- Recall the Primal Problem:

$$\min_w \max_{\alpha, \alpha_i \geq 0} L(w, \alpha)$$

- The Dual Problem:

$$\max_{\alpha, \alpha_i \geq 0} \min_w L(w, \alpha)$$

- **Theorem (weak duality):**

$$d^* = \max_{\alpha, \alpha_i \geq 0} \min_w L(w, \alpha) \leq \min_w \max_{\alpha, \alpha_i \geq 0} L(w, \alpha) = p^*$$

- **Theorem (strong duality):**

Iff there exist a saddle point of $L(w, \alpha)$

we have

$$d^* = p^*$$

Dual representation of the hard SVM QP

- We will start with the linearly separable case
- Instead of encoding the correct classification rule and constraint we will use Lagrange multipliers to encode it as part of the our minimization problem

$$\begin{aligned} & \text{Min } (\mathbf{w}^T \mathbf{w})/2 \\ & \text{s.t.} \\ & (\mathbf{w}^T \mathbf{x}_i + b) y_i \geq 1 \end{aligned}$$

Recall that Lagrange multipliers can be applied to turn the following problem:

$$L_{\text{primal}}(\mathbf{w}, b, \alpha) = \frac{1}{2} \mathbf{w} \cdot \mathbf{w} - \sum_{i=1}^N \alpha_i (y_i (\mathbf{w} \cdot \mathbf{x}_i + b) - 1)$$

The Dual Problem (Extra)

$$\max_{\alpha_i \geq 0} \min_{w,b} L(w,b,\alpha)$$

Dual formulation

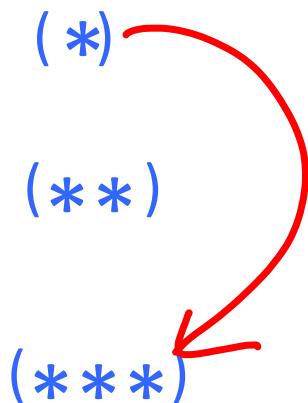
- We minimize L with respect to w and b first:

$$\nabla_w L(w,b,\alpha) = w - \sum_{i=1}^{\text{train}} \alpha_i y_i x_i = 0,$$

$$\nabla_b L(w,b,\alpha) = \sum_{i=1}^{\text{train}} \alpha_i y_i = 0,$$

Note that (*) implies:

$$w = \sum_{i=1}^{\text{train}} \alpha_i y_i x_i$$



- Plus (***) back to L , and using (**), we have:

$$\max_{\alpha \geq 0} L(w,b,\alpha) = \sum_{i=1} \alpha_i - \frac{1}{2} \sum_{i,j=1} \alpha_i \alpha_j y_i y_j (\mathbf{x}_i^T \mathbf{x}_j)$$

Summary: Dual for hard SVM (Extra)

Solving for \mathbf{w} that gives maximum margin:

1. Combine objective function and constraints into new objective function, using **Lagrange multipliers** α_i

$$L_{\text{primal}} = \frac{1}{2} \mathbf{w} \cdot \mathbf{w} - \sum_{i=1}^N \alpha_i (y_i (\mathbf{w} \cdot \mathbf{x}_i + b) - 1)$$

2. To minimize this **Lagrangian**, we take derivatives of \mathbf{w} and b and set them to 0:

Summary: Dual for hard SVM (Extra)

3. Substituting and rearranging gives the **dual** of the Lagrangian:

$$L_{dual} = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j \mathbf{x}_i \cdot \mathbf{x}_j$$

which we try to maximize (not minimize).

4. Once we have the α_i , we can substitute into previous equations to get \mathbf{w} and b .
5. This defines \mathbf{w} and b as **linear combinations of the training data**.

$$\mathbf{w} = \sum_{i=1}^{train} \alpha_i y_i \mathbf{x}_i$$

Summary: Dual SVM for linearly separable case

Dual formulation

$$\max_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$$

$$\sum_i \alpha_i y_i = 0$$

$$\alpha_i \geq 0 \quad \forall i$$

$n \times i$

$O(n)$
#para

w, b # para
 $O(p)$

$$\text{Min } (\underline{w^T w})/2$$

subject to the following inequality constraints:

For all x in class + 1

$$\underline{w^T x + b} \geq 1$$

For all x in class - 1

$$\underline{w^T x + b} \leq -1$$

}

A total of n constraints if we have n input samples



.

Easier than original QP, more efficient algorithms exist to find α_i ; e.g. SMO (see extra slides)

Dual formulation for linearly non-separable case

Dual target function:

$$\max_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{ij} \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$$

$$\sum_i \alpha_i y_i = 0$$

$$C \geq \alpha_i \geq 0, \forall i$$

Hyperparameter C
should be tuned
through k-folds CV

The only difference is that
the α are now
bounded

upper

$O(n)$
#para

prima linear soft
margin + Hinge
 $\sum \epsilon_i \alpha_i$
 w, b

$O(n + p + r)$

This is very similar to the
optimization problem in the linear
separable case, except that there
is an upper bound C on α_i now

Once again, efficient algorithm
exist to find α_i

Prediction via Dual Weights for linear case

Dual target function:

$$\max_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{ij} \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$$

$$\sum_i \alpha_i y_i = 0$$

$$C \geq \alpha_i \geq 0, \forall i$$

Hyperparameter C
should be tuned
through k-folds CV

The only difference is that
the α are now
bounded

$$f(x_{ts}) = \text{Sign}(w^T x_{ts} + b)$$
$$w = \sum_{i=1}^n \alpha_i y_i x_i$$

To evaluate a new sample x_{ts} we
need to compute:

$$w^T x_{ts} + b = \sum_{i \in \text{supportV}} \alpha_i y_i x_i^T x_{ts} + b$$

This is very similar to the
optimization problem in the linear
separable case, except that there is
an upper bound C on α_i now

Once again, efficient algorithm exist
to find α_i

Dual SVM – Training using Kernel Matrix

Our dual target function:

$$\max_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$$

$$\sum_i \alpha_i y_i = 0$$

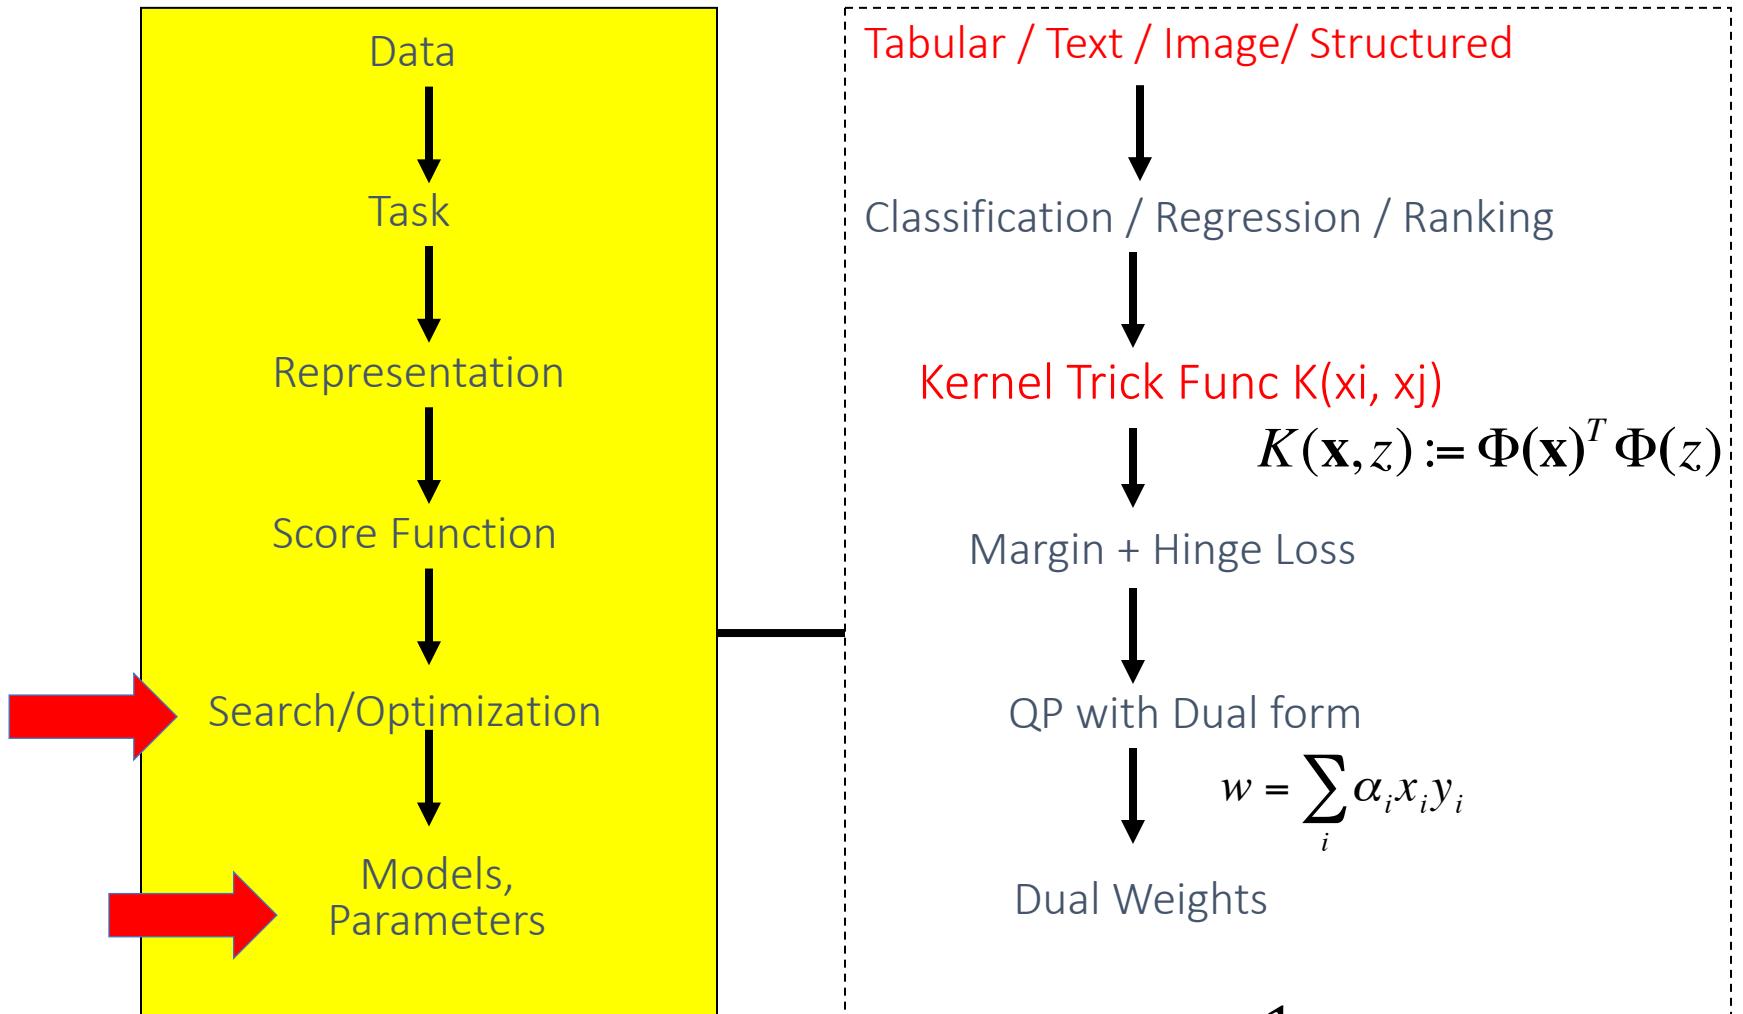
$$C \geq \alpha_i \geq 0, \forall i$$

Dot product among all training samples

$$\begin{matrix} & 1, 2, \dots, j, \dots, n \\ 1 & | \\ 2 & | \\ 3 & | \\ \vdots & | \\ i & | \\ \vdots & | \\ n & | \end{matrix}$$

matrix

This: Kernel Support Vector Machine



$$\underset{\mathbf{w}, b}{\operatorname{argmin}} \sum_{i=1}^p w_i^2 + C \sum_{i=1}^n \varepsilon_i$$

subject to $\forall \mathbf{x}_i \in D_{train} : y_i (\mathbf{x}_i \cdot \mathbf{w} + b) \geq 1 - \varepsilon_i$

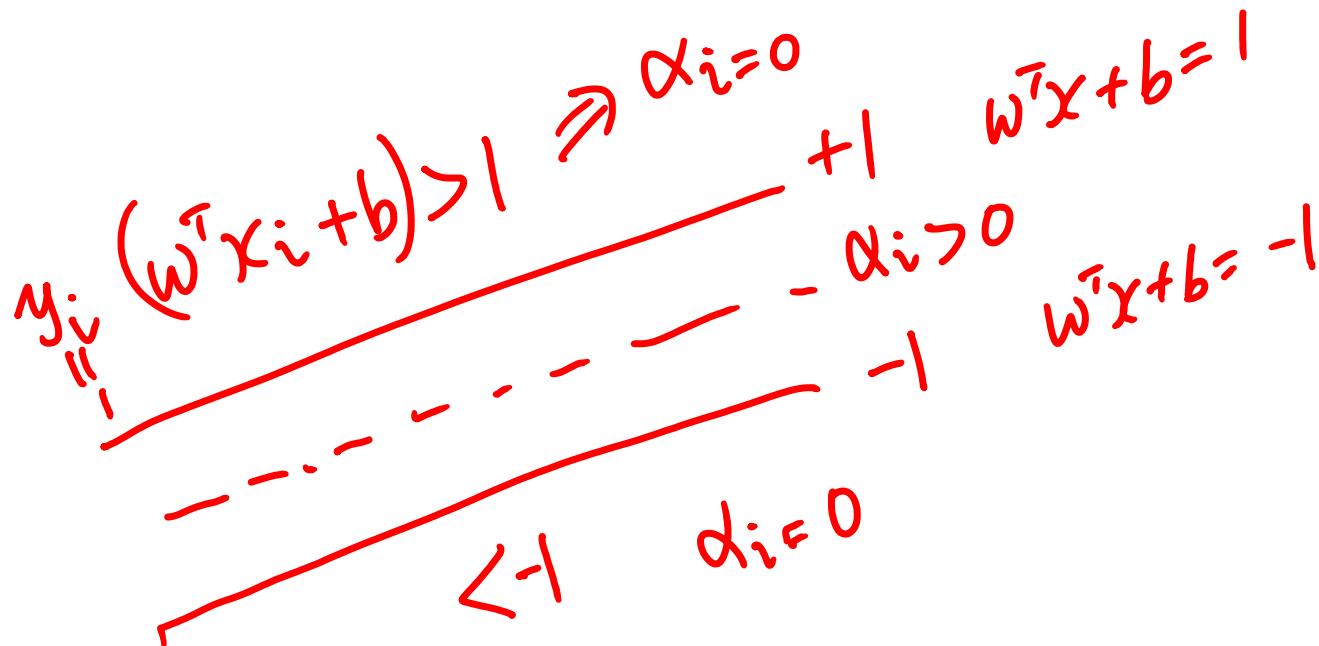
$$\max_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$$

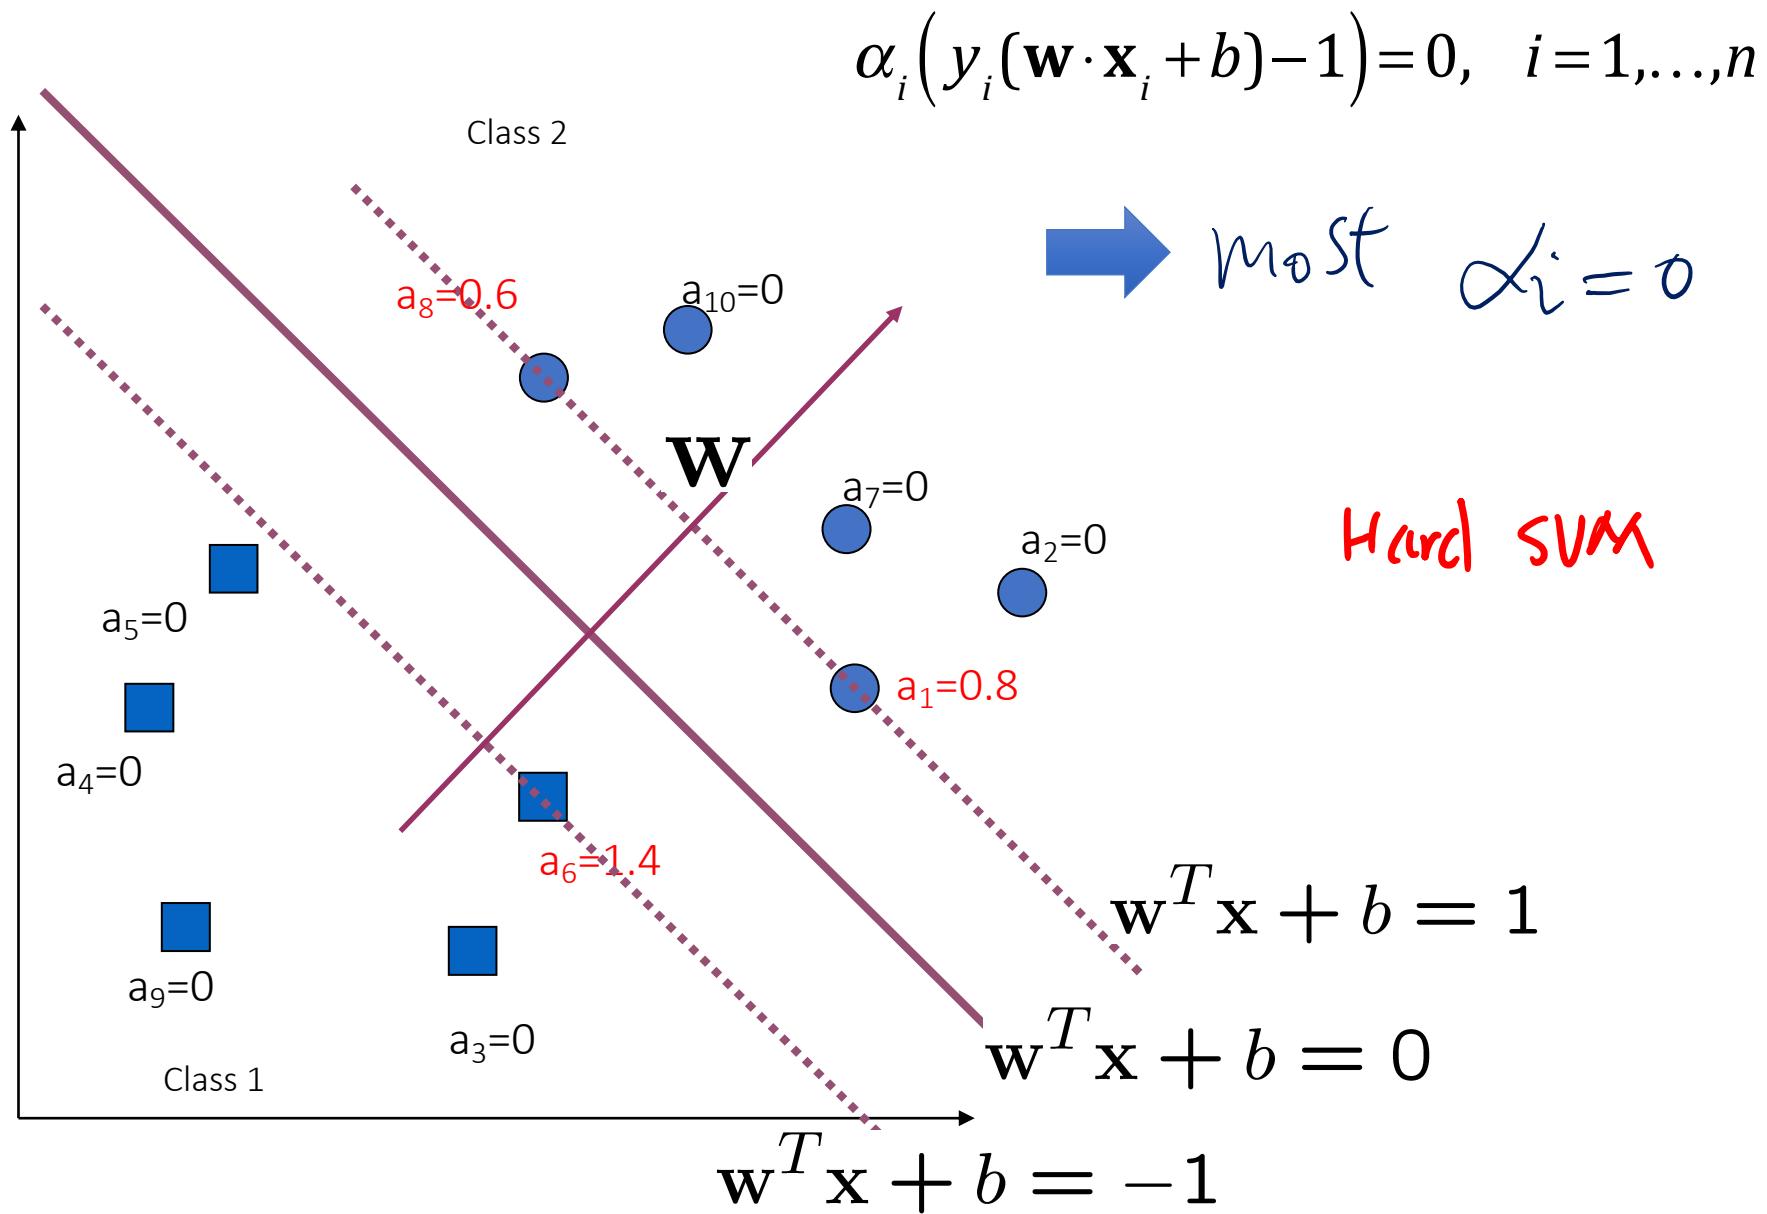
$$\sum_i \alpha_i y_i = 0, \quad \alpha_i \geq 0 \quad \forall i$$

Support vectors: non-zero a_i

- only a few a_i can be nonzero!!

$$\forall i \Rightarrow \alpha_i (y_i (\mathbf{w}^T \mathbf{x}_i + b) - 1) = 0, \quad i = 1, \dots, n$$

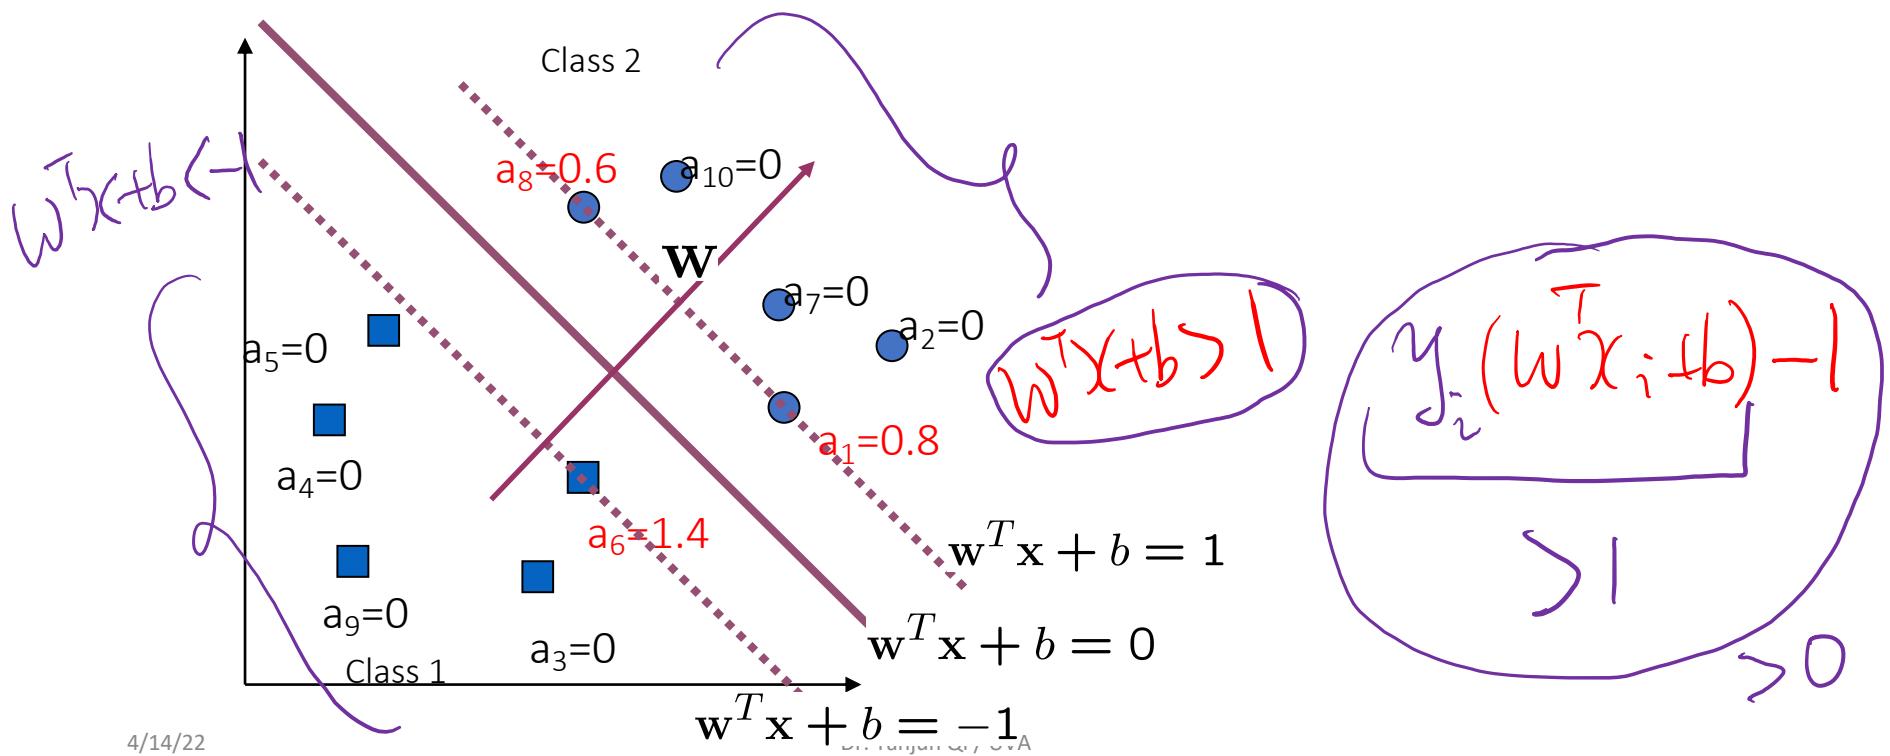




- only a few a_i can be nonzero! ie.

$$y_i(\mathbf{w} \cdot \mathbf{x}_i + b) = 1$$

$$\alpha_i(y_i(\mathbf{w} \cdot \mathbf{x}_i + b) - 1) = 0, \quad i=1,\dots,n \rightarrow \text{for most } \alpha_i \Rightarrow \alpha_i = 0$$



Support vectors: non-zero a_i

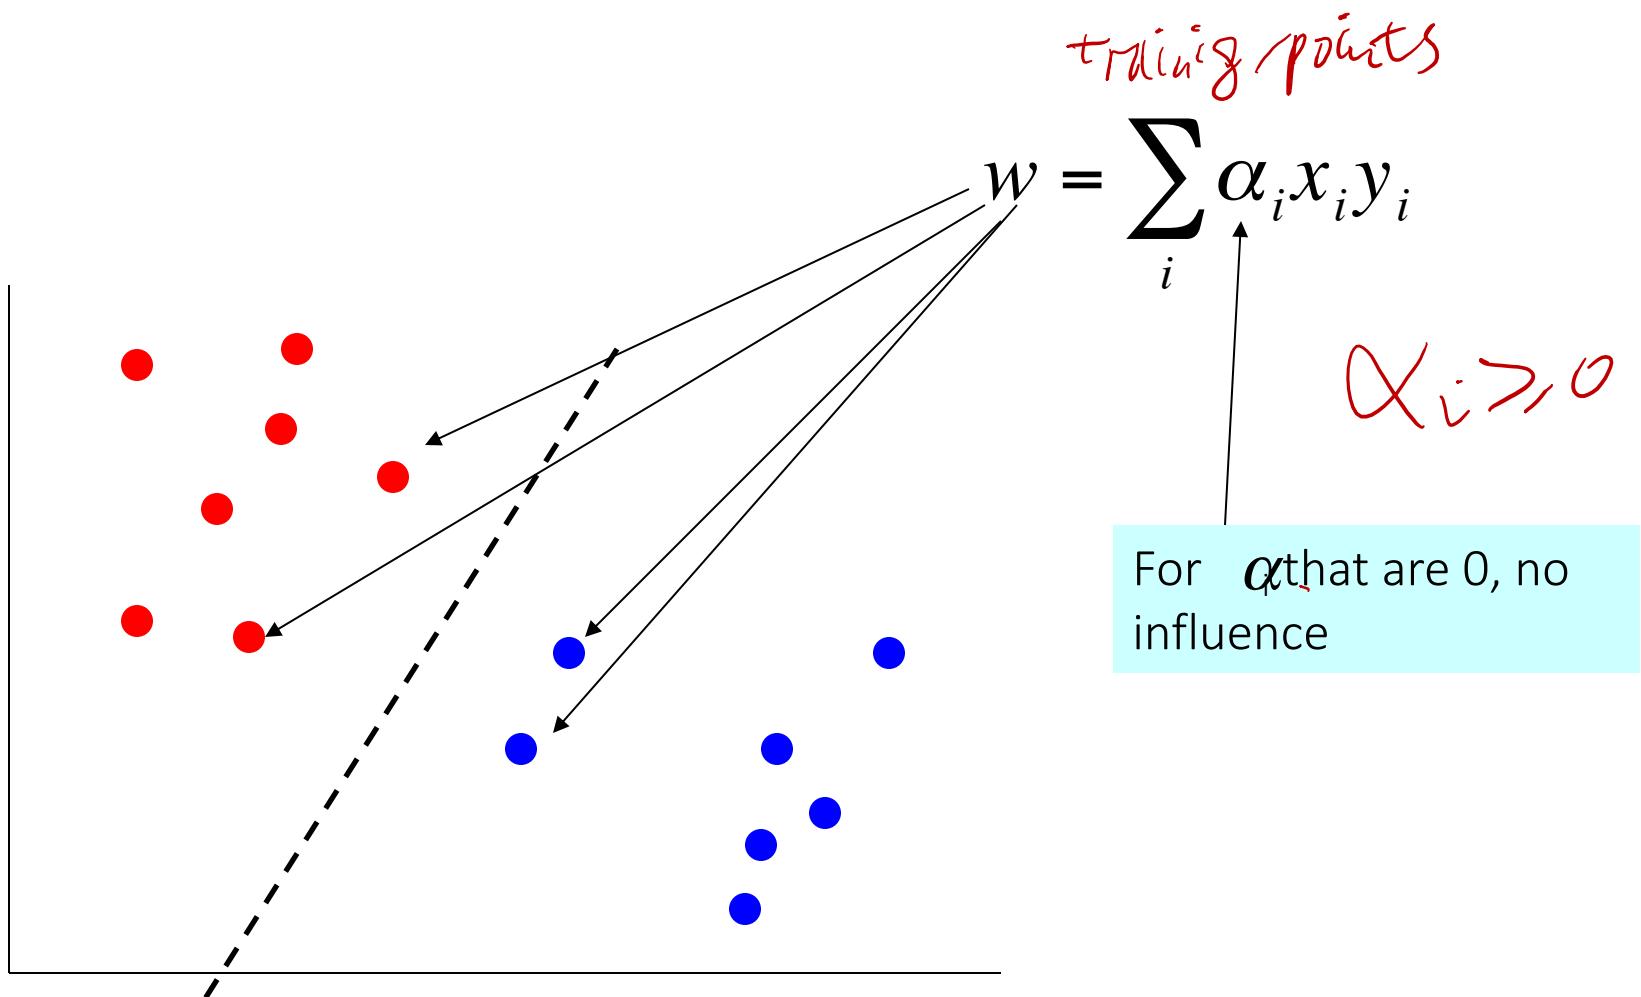
- only a few a_i can be nonzero!!

$$\alpha_i (y_i (\mathbf{w} \cdot \mathbf{x}_i + b) - 1) = 0, \quad i = 1, \dots, n$$

for most $\alpha_i \Rightarrow \alpha_i = 0$

We call the training data points whose a_i 's are nonzero the **support vectors (SV)**

Dual SVM - interpretation



Dual SVM– Testing

To evaluate a new sample x_{ts} we need to compute:

$$\hat{y}_{ts} = \text{sign}(w^T x_{ts} + b) = \text{sign}\left(\sum_{i=1..n} \alpha_i y_i x_i^T x_{ts} + b\right)$$

$$\rightarrow \hat{y}_{ts} = \text{sign}\left(\sum_{i \in \text{SupportVectors}} \alpha_i y_i (x_i^T x_{ts}) + b\right)$$

\downarrow
 $\alpha_i > 0$

For α_i that are 0,
no influence

Dot product with (“all” ??)
training samples

Support Vectors for the Soft-SVM

$$\alpha_i (y_i (\mathbf{w} \cdot \mathbf{x}_i + b) - 1 + \varepsilon_i) = 0, \quad i = 1, \dots, n$$

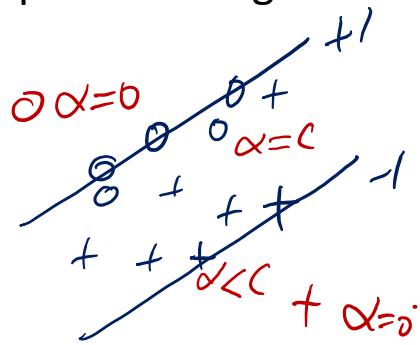
- Support vectors are

- Samples on the margin:

$$y_i (\mathbf{x}_i \cdot \mathbf{w} + b) = 1,$$

$$0 < \alpha_i < C$$

- Samples violating the margin (mostly inside the margin area):



$$y_i (\mathbf{x}_i \cdot \mathbf{w} + b) < 1,$$

$$\alpha_i = C$$

More in L11Extra-SVMoptimDual

Value C and Number of Support Vectors (no clear relation!!!)

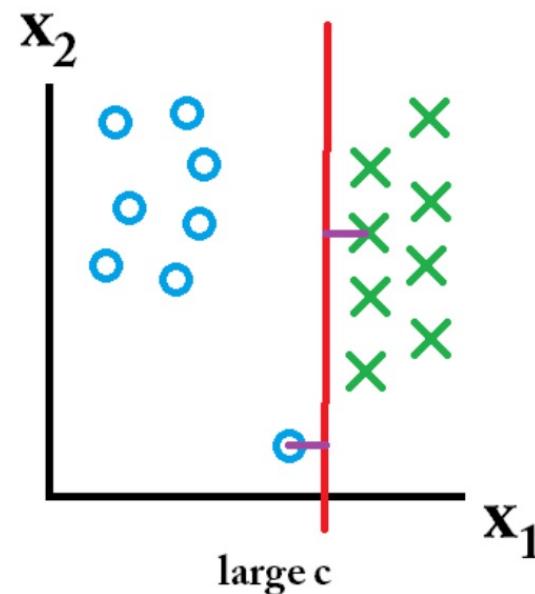
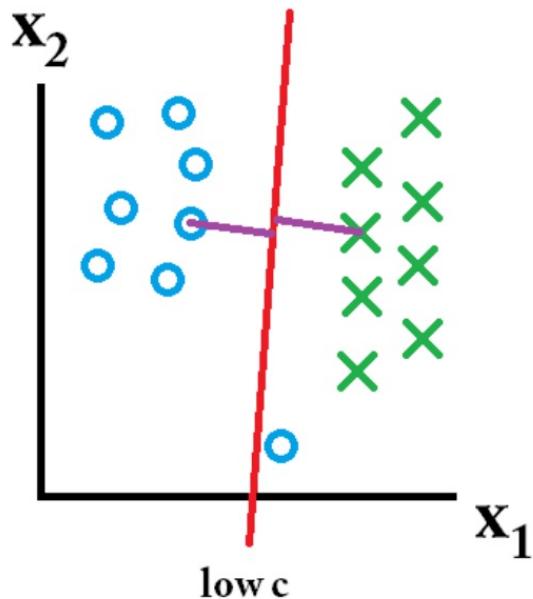
C is how hard we want to punish misclassified examples in the training set.

So for large values of C, it would make sense that if misclassified examples are severely punished, then it will choose a small margin with not many support vectors.

For small values of C, it would broaden the margin, and as a result, end up getting more points inside of the margin (if not easily separable), so there would be possibly more support vectors.

I thought that this [StackExchange post](#) described it pretty well.

As in this example, the small c has more "support vectors" than the large c



<https://github.com/scikit-learn/scikit-learn/issues/7955>

<https://stats.stackexchange.com/questions/31066/what-is-the-influence-of-c-in-svms-with-linear-kernel>

Why do SVMs work?

$x \rightarrow \phi(x)$ e.g. RBF

- If we are using huge features spaces (e.g., with kernels), how come we are not overfitting the data?
 - ✓ Number of parameters remains the same (and most are set to 0) $O(n)$, α_i $i=1, \dots, n$
 - ✓ While we have a lot of inputs, at the end we only care about the support vectors and these are usually a small group of samples
 - ✓ The maximizing of the margin acts as a sort of regularization term leading to reduced overfitting

Summary of SVM

❑ Support Vector Machine (SVM)

- ✓ History of SVM
- ✓ Large Margin Linear Classifier
- ✓ Define Margin (M) in terms of model parameter
- ✓ Optimization to learn model parameters (w, b)
- ✓ Non linearly separable case
- ✓ Optimization with dual form
- ✓ Nonlinear decision boundary
- ✓ Practical Guide
 - ✓ File format / LIBSVM
 - ✓ Feature preprocessing
 - ✓ Model selection
 - ✓ Pipeline procedure

(Recap)

kNN : $\hat{y}_{ts} = \frac{1}{k} \sum_{i \in k \text{ Neighbors of } x_{ts}} y_i$
find k neighbor of x_{ts} $\sim O(n^k)$

SVM : $\hat{y}_{ts} = \sum_{i \in SV} \alpha_i y_i k(\vec{x}_i, \vec{x}_{ts}) + b$
 $\# \text{para} \sim O(n)$

Logistic Regression / Linear Classifier
 $\hat{y}_{ts} = \sigma(W^T x_{ts} + b)$
 $\# \text{para} \sim O(p)$

References

- Big thanks to Prof. Ziv Bar-Joseph and Prof. Eric Xing @ CMU for allowing me to reuse some of his slides
- Elements of Statistical Learning, by Hastie, Tibshirani and Friedman
- Prof. Andrew Moore @ CMU's slides
- Tutorial slides from Dr. Tie-Yan Liu, MSR Asia
- A Practical Guide to Support Vector Classification Chih-Wei Hsu, Chih-Chung Chang, and Chih-Jen Lin, 2003-2010
- Tutorial slides from Stanford “Convex Optimization I — Boyd & Vandenberghe