

UVA CS 4774: Machine Learning

S5: Lecture 26: Reinforcement Learning

Dr. Yanjun Qi

University of Virginia
Department of Computer Science

Course Content Plan → Regarding Tasks

~~Regression (supervised)~~

Y is a continuous

~~Learning theory~~

About $f()$

~~Classification (supervised)~~

Y is a discrete

~~Unsupervised models~~

NO Y

~~Graphical models~~

About interactions among Y, X₁,.., X_p






Reinforcement Learning


Learn to Interact with environment

Outline

- Examples of RL applications
- Defining an RL problem
 - Markov Decision Processes
- Solving an RL problem
 - Dynamic Programming
 - Monte Carlo methods
 - Temporal-Difference learning

Where Machine Learning is being used or can be useful?





Classes of Learning Problems

Supervised Learning:

Data: (x, y)

x is data, y is label

Goal: Learn function
to map $x \rightarrow y$

Example:



This thing is an apple.

Unsupervised Learning:

Data: x

x is data, no labels!

Goal: Learn underlying
structure

Example:



This thing is like
the other thing.

Reinforcement Learning:


Data: state-action pairs

Goal: Maximize future
rewards over many steps

Example:




Eat this thing because it
will keep you alive.




The Machine Learning Stack!

What can be learned?



Sensors



Lidar



Camera
(Visible, Infrared)



Radar



GPS




Stereo Camera




Microphone




Networking
(Wired, Wireless)




IMU



Representations






Knowledge / Reasoning

Image Recognition:
If it looks like a duck




Audio Recognition:
Quacks like a duck




Activity Recognition:
Swims like a duck






Actions




The Full Stack



Reinforcement Learning

- Learning to interact with an environment
 - Robots, games, process control
 - With limited human training
 - Where the ‘right thing’ isn’t obvious
- Supervised Learning:
 - Goal: $f(x) = y$
 - Data: $[< x_1, y_1 >, \dots, < x_n, y_n >]$
- Reinforcement Learning:
 - Goal:
$$\text{Maximize } \sum_{i=1}^{\infty} \text{Reward}(State_i, Action_i)$$
 - Data:
$$Reward_i, State_{i+1} = \text{Interact}(State_i, Action_i)$$



credit: Geoff Hulten


History of Reinforcement Learning

- Roots in the psychology of animal learning (Thorndike, 1911).
- Another independent thread was the problem of optimal control, and its solution using dynamic programming (Bellman, 1957).
- Idea of temporal difference learning (on-line method), e.g., playing board games (Samuel, 1959).
- A major breakthrough was the discovery of Q-learning (Watkins, 1989).

A Success Story


- TD Gammon (Tesauro, G., 1992)
 - A Backgammon playing program.
 - Application of temporal difference learning.
 - The basic learner is a neural network.
 - It trained itself to the world class level by playing against itself and learning from the outcome. So smart!!
 - More information:

<http://www.research.ibm.com/massive/tdl.html>



TD-Gammon – Tesauro ~1995

State: Board State
Actions: Valid Moves
Reward: Win or Lose



- Net with 80 hidden units, initialize to random weights
- Select move based on network estimate & shallow search
- Learn by playing against itself
- 1.5 million games of training
-> competitive with world class players


credit: Geoff Hulten

Examples of Reinforcement Learning

- How should a robot behave so as to optimize its “performance”? **(Robotics)**
- How to automate the motion of a helicopter? **(Control Theory)**
- How to make a good chess-playing program? **(Artificial Intelligence)**





Resource allocation in datacenters



- A Hybrid Reinforcement Learning Approach to Autonomic Resource Allocation
 - Tesauro, Jong, Das, Bennani (IBM)
 - ICAC 2006

Atari 2600 games



- Same model/parameters for ~50 games

credit: Geoff Hulten

<https://storage.googleapis.com/deepmind-media/dqn/DQNNaturePaper.pdf>

Robotics and Locomotion





Figure 5: Time-lapse images of a representative *Quadruped* policy traversing gaps (left); and navigating obstacles (right)

State:
Joint States/Velocities
Accelerometer/Gyroscope
Terrain
Actions: Apply Torque to Joints
Reward: Velocity – { stuff }



https://youtu.be/hx_bgoTF7bs

credit: Geoff Hulten



2017 paper <https://arxiv.org/pdf/1707.02286.pdf>

Alpha Go

- Learning how to beat humans at ‘hard’ games (search space too big)
- Far surpasses (Human) Supervised learning
- Algorithm learned to outplay humans at chess in 24 hours





State: Board State
Actions: Valid Moves
Reward: Win or Lose






credit: Geoff Hulten

Deep Reinforcement Learning


- Human
- So what's DEEP RL?





AlphaGO: Learning Pipeline

- Combine Supervised Learning (SL) and RL to learn the search direction in Monte Carlo Tree Search



- SL policy Network
 - Prior search probability or potential
- Rollout:
 - combine with MCTS for quick simulation on leaf node
- Value Network:
 - Build the Global feeling on the leaf node situation

AlphaGo {Fan, Lee, Master} × AlphaGo Zero:

- supervised learning from human expert positions × from scratch by self-play reinforcement learning (“tabula rasa”)
- additional (auxiliary) input features × only the black and white stones from the board as input features
- separate policy and value networks × single neural network
- tree search using also Monte Carlo rollouts × simpler tree search using only the single neural network to both evaluate positions and sample moves
- (AlphaGo Lee) distributed machines + 48 tensor processing units (TPUs) × single machines + 4 TPUs
- (AlphaGo Lee) several months of training time × 72 h of training time (outperforming AlphaGo Lee after 36 h)



USER DOCUMENTATION

- Introduction
- Installation
- Algorithms
- Running Experiments
- Experiment Outputs
- Plotting Results

INTRODUCTION TO RL

- Part 1: Key Concepts in RL
- Part 2: Kinds of RL Algorithms
- Part 3: Intro to Policy Optimization

RESOURCES

- Spinning Up as a Deep RL Researcher
- Key Papers in Deep RL

Benchmarks for Spinning Up Implementations

Table of Contents

- [Benchmarks for Spinning Up Implementations](#)
 - [Performance in Each Environment](#)
 - [HalfCheetah: PyTorch Versions](#)
 - [HalfCheetah: Tensorflow Versions](#)
 - [Hopper: PyTorch Versions](#)
 - [Hopper: Tensorflow Versions](#)
 - [Walker2d: PyTorch Versions](#)
 - [Walker2d: Tensorflow Versions](#)
 - [Swimmer: PyTorch Versions](#)
 - [Swimmer: Tensorflow Versions](#)
 - [Ant: PyTorch Versions](#)
 - [Ant: Tensorflow Versions](#)
 - [Experiment Details](#)
 - [PyTorch vs Tensorflow](#)

We benchmarked the Spinning Up algorithm implementations in five environments from the [MuJoCo](#) Gym task suite: HalfCheetah, Hopper, Walker2d, Swimmer, and Ant.

Performance in Each Environment

What is special about RL?

- RL is learning how to map states to actions, so as to **maximize** a numerical **reward** over time.
- Unlike other forms of learning, it is a multistage decision-making process (often **Markovian**).
- An RL agent learn by **trial-and-error**. (Not entirely supervised, but interactive)
- Actions may affect not only the immediate reward but also subsequent rewards (**Delayed effect**).

Outline

- Examples of RL applications
- Defining an RL problem
 - Markov Decision Processes
- Solving an RL problem
 - Dynamic Programming
 - Monte Carlo methods
 - Temporal-Difference learning

Elements of RL

- A **policy**
 - A map from **state space** to **action space**.
 - May be stochastic.
- A **reward function**
 - It maps each state (or, state-action pair) to a real number, called **reward**.
- A **value function**
 - Value of a state (or, state-action pair) is the **total expected reward**, starting from that state (or, state-action pair).

Setup for Reinforcement Learning

Markov Decision Process (environment)

- Discrete-time stochastic control process
- Each time step, s :
 - Agent chooses action a from set A_s
 - Moves to new state with probability:
 - $P_a(s, s')$
 - Receives reward: $R_a(s, s')$
- Every outcome depends on s and a
 - Nothing depends on previous states/actions

Policy

(agent's behavior)

- $\pi(s)$ – The action to take in state s
- Goal maximize: $\sum_{t=0}^{\infty} \gamma^t R_{a_t}(s_t, s_{t+1})$
 - $a_t = \pi(s_t)$
 - $0 \leq \gamma < 1$ – Tradeoff immediate vs future

$$V^{\pi}(s) = \sum_{s'} P_{\pi(s)}(s, s') * (R_{\pi(s)}(s, s') + \gamma V^{\pi}(s'))$$

Reward for making that move
credit: Geoff Hulten

Probability of moving to each state
Value of being in that state

Simple Example of Agent in an Environment

State:

Map Locations

$\{<0,0>, <1,0> \dots <3,3>\}$

Actions:

Move within map

Reaching chest ends episode


$$A_{0,0} = \{east, south\}$$

$$A_{1,0} = \{east, south, west\}$$

$$A_{2,0} = \{\phi\}$$

...

$$A_{2,2} = \{north, west\}$$



Reward:

100 at chest

0 for others

$$R_{east}(<1,0>, <2,0>) = 100$$

$$R_{north}(<2,1>, <2,0>) = 100$$

$$R_*(*,*) = 0$$



credit: Geoff Hulten


Policies

Policy

$$\pi(s) = a$$

- $\pi(<0,0>) = \{ south \}$
- $\pi(<0,1>) = \{ east \}$
- $\pi(<0,2>) = \{ east \}$
- $\pi(<1,0>) = \{ east \}$
- $\pi(<1,1>) = \{ north \}$
- $\pi(<1,2>) = \{ north \}$
- $\pi(<2,0>) = \{ \phi \}$
- $\pi(<2,1>) = \{ west \}$
- $\pi(<2,2>) = \{ north \}$

Policy could be better



$R_{east}(<1,0>, <2,0>)$	= 100
$R_{north}(<2,1>, <2,0>)$	= 100
$R_*(*,*)$	= 0
$\gamma = 0.5$	

Evaluating Policies

$$V^\pi(s) = \sum_{i=0}^{\infty} \gamma^i r_{i+1}$$


$$V^\pi(<1,0>) = \gamma^0 * 100$$

$$V^\pi(<1,1>) = \gamma^0 * 0 + \gamma^1 * 100$$

Move to <1,0>
Move to <1,1>
Move to <1,0>
Move to <2,0>

$$V^\pi(<0,0>) = \boxed{\gamma^0 * 0} + \boxed{\gamma^1 * 0} + \boxed{\gamma^2 * 0} + \boxed{\gamma^3 * 100}$$


Robot in a room



actions: UP, DOWN, LEFT, RIGHT

UP


80% move UP
10% move LEFT
10% move RIGHT



- reward +1 at [4,3], -1 at [4,2]
- reward -0.04 for each step
- what's the strategy to achieve max reward?
- what if the actions were deterministic?

Other examples

- pole-balancing
 - TD-Gammon [Gerry Tesauro]
 - helicopter [Andrew Ng]
-
- no teacher who would say “good” or “bad”
 - is reward “10” good or bad?
 - rewards could be delayed
 - similar to control theory
 - more general, fewer constraints
 - explore the environment and learn from experience
 - not just blind search, try to be smart about it



How Reinforcement Learning is Different

- Delayed Reward
- Agent chooses training data
- Explore vs Exploit (Life long learning)
- Very different terminology (can be confusing)

credit: Geoff Hulten

Outline

- Examples of RL applications
- Defining an RL problem
 - Markov Decision Processes
- Solving an RL problem
 - Dynamic Programming
 - Monte Carlo methods
 - Temporal-Difference learning

The Precise Goal / Popular RL Algorithms

- To find a **policy** that maximizes the **Value function**.
 - transitions and rewards usually not available
- There are different approaches to achieve this goal in various situations.
- Value iteration and Policy iteration are two more classic approaches to this problem. But essentially both are **dynamic programming**.
- Q-learning is a more recent approaches to this problem. Essentially it is a **temporal-difference method**.

(1) Dynamic programming

- main idea
 - use value functions to structure the search for good policies
 - need a perfect model of the environment
- two main components
 - policy evaluation: compute V^π from π
 - policy improvement: improve π based on V^π
- start with an arbitrary policy
- repeat evaluation/improvement until convergence


Value functions

- state value function: $V^\pi(s)$
 - expected return when starting in s and following π
- state-action value function: Q-function: $Q^\pi(s,a)$
 - expected return when starting in s , performing a , and following π

- useful for finding the optimal policy
 - can estimate from experience
 - pick the best action using $Q^\pi(s,a)$

$$V^\pi(s) = \sum_a \pi(s, a) \sum_{s'} P_{ss'}^a [r_{ss'}^a + \gamma V^\pi(s')] = \sum_a \pi(s, a) Q^\pi(s, a)$$

- Bellman equation



Using DP

- need complete model of the environment and rewards
 - robot in a room
 - state space, action space, transition model
- can we use DP to solve
 - robot in a room?
 - back gammon?
 - helicopter?

Outline

- Examples of RL applications
- Defining an RL problem
 - Markov Decision Processes
- Solving an RL problem
 - Dynamic Programming
 - Monte Carlo methods
 - Temporal-Difference learning

Monte Carlo methods

- don't need full knowledge of environment
 - just experience, or
 - simulated experience
- but similar to DP
 - policy evaluation, policy improvement
- averaging sample returns
 - defined only for episodic tasks


Computing return from rewards

- episodic (vs. continuing) tasks
 - “game over” after N steps
 - optimal policy depends on N; harder to analyze
- additive rewards
 - $V(s_0, s_1, \dots) = r(s_0) + r(s_1) + r(s_2) + \dots$
 - infinite value for continuing tasks
- discounted rewards
 - $V(s_0, s_1, \dots) = r(s_0) + \gamma * r(s_1) + \gamma^2 * r(s_2) + \dots$
 - value bounded if rewards bounded

Monte Carlo policy evaluation

- want to estimate $V^\pi(s)$
 - = expected return starting from s and following π
- estimate as average of observed returns in state s

- first-visit MC
 - average returns following the first visit to state s



$$V^\pi(s) \approx (2 + 1 - 5 + 4)/4 = 0.5$$

Maintaining exploration

- deterministic/greedy policy won't explore all actions
 - don't know anything about the environment at the beginning
 - need to try all actions to find the optimal one
- maintain exploration
 - use *soft* policies instead: $\pi(s,a) > 0$ (for all s,a)
- ϵ -greedy policy
 - with probability $1-\epsilon$ perform the optimal/greedy action
 - with probability ϵ perform a random action
 - will keep exploring the environment
 - slowly move it towards greedy policy: $\epsilon \rightarrow 0$

Simulated experience

- 5-card draw poker
 - $s_0: A\clubsuit, A\heartsuit, 6\spades, A\heartsuit, 2\spades$
 - $a_0: \text{discard } 6\spades, 2\spades$
 - $s_1: A\clubsuit, A\heartsuit, A\heartsuit, A\spades, 9\spades + \text{dealer takes 4 cards}$
 - return: +1 (probably)
- DP
 - list all states, actions, compute $P(s,a,s')$
 - $P([A\clubsuit, A\heartsuit, 6\spades, A\heartsuit, 2\spades], [6\spades, 2\spades], [A\spades, 9\spades, 4]) = 0.00192$
- MC
 - all you need are sample episodes
 - let MC play against a random policy, or itself, or another algorithm

Summary of Monte Carlo

- don't need model of environment
 - averaging of sample returns
 - only for episodic tasks
- learn from sample episodes or simulated experience
- can concentrate on “important” states
 - don't need a full sweep
- need to maintain exploration
 - use soft policies

Outline

- Examples of RL applications
- Defining an RL problem
 - Markov Decision Processes
- Solving an RL problem
 - Dynamic Programming
 - Monte Carlo methods
 - Temporal-Difference learning

Temporal Difference Learning


- combines ideas from MC and DP
 - like MC: learn directly from experience (don't need a model)
 - like DP: learn from values of successors
 - works for continuous tasks, usually faster than MC
- constant-alpha MC:
 - have to wait until the end of episode to update

$$V(s_t) \leftarrow V(s_t) + \alpha [R_t - V(s_t)]$$

target

- simplest TD
 - update after every step, based on the successor

$$V(s_t) \leftarrow V(s_t) + \alpha [r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$




Value functions

- state value function: $V^\pi(s)$
 - expected return when starting in s and following π
- state-action value function: Q-function: $Q^\pi(s,a)$
 - expected return when starting in s , performing a , and following π

- useful for finding the optimal policy
 - can estimate from experience
 - pick the best action using $Q^\pi(s,a)$

$$V^\pi(s) = \sum_a \pi(s, a) \sum_{s'} P_{ss'}^a [r_{ss'}^a + \gamma V^\pi(s')] = \sum_a \pi(s, a) Q^\pi(s, a)$$

- Bellman equation



Optimal value functions

- there's a set of *optimal* policies
 - V^π defines partial ordering on policies
 - they share the same optimal value function


$$V^*(s) = \max_{\pi} V^\pi(s)$$

- Bellman optimality equation

$$V^*(s) = \max_a \sum_{s'} P_{ss'}^a [r_{ss'}^a + \gamma V^*(s')]$$

- system of n non-linear equations
 - solve for $V^*(s)$
 - easy to extract the optimal policy
- having $Q^*(s,a)$ makes it even simpler

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$



credit: Peter Bodí

Q-learning

- before: on-policy algorithms
 - start with a random policy, iteratively improve
 - converge to optimal

- Q-learning: off-policy
 - use any policy to estimate Q

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right]$$


- Q directly approximates Q^* (Bellman optimality eqn)
- independent of the policy being followed
- only requirement: keep updating each (s,a) pair

- Sarsa

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t) \right]$$

Sarsa


- again, need $Q(s,a)$, not just $V(s)$



$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

- control
 - start with a random policy
 - update Q and π after each step
 - again, need ϵ -soft policies

The RL Intro book



Richard Sutton, Andrew Barto
Reinforcement Learning,
An Introduction

<http://www.cs.ualberta.ca/~sutton/book/the-book.html>

credit: Peter Bodí

Summary

Reinforcement Learning:

- Goal: Maximize $\sum_{i=1}^{\infty} \text{Reward}(\text{State}_i, \text{Action}_i)$
- Data: $\text{Reward}_{i+1}, \text{State}_{i+1} = \text{Interact}(\text{State}_i, \text{Action}_i)$



Many (awesome) recent successes:

- Robotics
- Surpassing humans at difficult games
- Doing it with (essentially) zero human knowledge

Challenges:

- When the episode can end without reward
- When there is a ‘narrow’ path to reward
- When there are many states and actions

(Simple) Approaches:

- Q-Learning $\hat{Q}(s, a) \rightarrow$ discounted reward of action
- Policy Gradients \rightarrow Probability distribution over A_s
- Reward Shaping
- Memory
- Lots of parameter tweaking...

- Key Papers in Deep RL
 - 1. Model-Free RL
 - 2. Exploration
 - 3. Transfer and Multitask RL
 - 4. Hierarchy
 - 5. Memory
 - 6. Model-Based RL
 - 7. Meta-RL
 - 8. Scaling RL
 - 9. RL in the Real World
 - 10. Safety
 - 11. Imitation Learning and Inverse Reinforcement Learning
 - 12. Reproducibility, Analysis, and Critique
 - 13. Bonus: Classic Papers in RL Theory or Review

References


- RL slides from Rich Nguven
- RL Slides from Geoff Hulten
- RL slides from Eric Xing
- RL slides from Peter Bodik

credit: Geoff Hulten

- Vanilla Policy Gradient
 - Background
 - Documentation
 - References
- Trust Region Policy Optimization
 - Background
 - Documentation
 - References
- Proximal Policy Optimization
 - Background
 - Documentation
 - References
- Deep Deterministic Policy Gradient
 - Background
 - Documentation
 - References
- Twin Delayed DDPG
 - Background
 - Documentation
 - References
- Soft Actor-Critic
 - Background
 - Documentation
 - References


Gym – toolkit for reinforcement learning

CartPole



Reward +1 per step the pole remains up

MountainCar



Reward 200 at flag -1 per step

```
import gym

env = gym.make('CartPole-v0')

import random
import QLearning # Your implementation goes here...
import Assignment7Support

trainingIterations = 20000

qlearner = QLearning.QLearning(<Parameters>)

for trialNumber in range(trainingIterations):
    observation = env.reset()
    reward = 0
    for i in range(300):
        env.render() # Comment out to make much faster...

        currentState = ObservationToStateSpace(observation)
        action = qlearner.GetAction(currentState, <Parameters>)

        oldState = ObservationToStateSpace(observation)
        observation, reward, isDone, info = env.step(action)
        newState = ObservationToStateSpace(observation)

        qlearner.ObserveAction(oldState, action, newState, reward, ...)

        if isDone:
            if(trialNumber%1000) == 0:
                print(trialNumber, i, reward)
            break

# Now you have a policy in qlearner - use it...
```

Q learning

Learn a policy $\pi(s)$ that optimizes $V^\pi(s)$ for all states, using:

- No prior knowledge of state transition probabilities: $P_a(s, s')$
- No prior knowledge of the reward function: $R_a(s, s')$

Approach:


- Initialize estimate of discounted reward for every state/action pair: $\hat{Q}(s, a) = 0$
- Repeat (for a while):
 - Take a random action a from A_s
 - Receive s' and $R_a(s, s')$ from environment
 - Update $\hat{Q}(s, a) = R_a(s, s') + \gamma \max_{a'} \hat{Q}(s', a')$
 - Random restart if in terminal state

$$\alpha_v = \frac{1}{1 + \text{visits}(s, a)}$$

$$\text{Exploration Policy: } P(a_i, s) = \frac{k^{\hat{Q}(s, a_i)}}{\sum_j k^{\hat{Q}(s, a_j)}}$$

Example of Q learning (round 1)

- Initialize \hat{Q} to 0
- Random initial state = $< 1,1 >$
- Random action from $A_{<1,1>} = east$
 - $s' = < 2,1 >$
 - $R_a(s, s') = 0$
- Update $\hat{Q}(< 1,1 >, east) = 0$
- Random action from $A_{<2,1>} = north$
 - $s' = < 2,0 >$
 - $R_a(s, s') = 100$
- Update $\hat{Q}(< 2,1 >, north) = 100$
- No more moves possible, start again...




$$\hat{Q}(s, a) = R_a(s, s') + \gamma \max_{a'} \hat{Q}_{n-1}(s', a')$$

credit: Peter Bodí

Example of Q learning (round 2)

- Round 2: Random initial state = $< 2,2 >$
- Random action from $A_{<2,2>} = north$
 - $s' = < 2,1 >$
 - $R_a(s, s') = 0$
- Update $\hat{Q}(< 2,1 >, north) = 0 + \gamma * 100$
- Random action from $A_{<2,1>} = north$
 - $s' = < 2,0 >$
 - $R_a(s, s') = 100$
- Update $\hat{Q}(< 2,1 >, north) = still 100$
- No more moves possible, start again...



$$\hat{Q}(s, a) = R_a(s, s') + \gamma \max_{a'} \hat{Q}_{n-1}(s', a')$$


$$\gamma = 0.5$$

credit: Peter Bodí

Example of Q learning (some acceleration...)

$$\hat{Q}(s, a) = R_a(s, s') + \gamma \max_{a'} \hat{Q}_{n-1}(s', a')$$
$$\gamma = 0.5$$


- Random Initial State $< 0,0 >$
- Update $\hat{Q}(< 1,1 >, east) = 50$
- Update $\hat{Q}(< 1,2 >, east) = 25$



Example of Q learning (some acceleration...)

$$\hat{Q}(s, a) = R_a(s, s') + \gamma \max_{a'} \hat{Q}_{n-1}(s', a')$$
$$\gamma = 0.5$$


- Random Initial State $< 0,2 >$
- Update $\hat{Q}(< 0,1 >, east) = 25$
- Update $\hat{Q}(< 1,0 >, east) = 100$



Example of Q learning (\hat{Q} after many, many runs...)


- \hat{Q} converged
- Policy is:

$$\pi(s) = \operatorname{argmax}_{a \in A_s} \hat{Q}(s, a)$$




Challenges for Reinforcement Learning

- When there are many states and actions
- When the episode can end without reward
- When there is a ‘narrow’ path to reward




Memory

- Retrain on previous explorations
 - Maintain samples of:
 $P_a(s, s')$
 $R_a(s, s')$
- Useful when
 - It is cheaper to use some RAM/CPU than to run more simulations
 - It is hard to get reward so you want to leverage it for as much as possible when it happens



Robot in a room




- states
- actions
- rewards
- what is the solution?

actions: UP, DOWN, LEFT, RIGHT


UP

80% move UP
10% move LEFT
10% move RIGHT



reward +1 at [4,3], -1 at [4,2]
reward -0.04 for each step


Is this a solution?



- only if actions deterministic
 - not in this case (actions are stochastic)
- solution/policy
 - mapping from each state to an action


State representation

- pole-balancing
 - move car left/right to keep the pole balanced
- state representation
 - position and velocity of car
 - angle and angular velocity of pole
- what about *Markov property*?
 - would need more info
 - noise in sensors, temperature, bending of pole
- solution
 - coarse discretization of 4 state variables
 - left, center, right
 - totally non-Markov, but still works




Splitting and aggregation

- want to discretize the state space
 - learn the best discretization during training
- splitting of state space
 - start with a single state
 - split a state when different *parts* of *that state* have different values



- state aggregation
 - start with many states
 - merge states with similar values



credit: Peter Bodí

Designing rewards

- robot in a maze
 - episodic task, not discounted, +1 when out, 0 for each step
- chess
 - GOOD: +1 for winning, -1 losing
 - BAD: +0.25 for taking opponent's pieces
 - high reward even when lose
- rewards
 - rewards indicate what we want to accomplish
 - NOT how we want to accomplish it
- shaping
 - positive reward often very “far away”
 - rewards for achieving subgoals (domain knowledge)
 - also: adjust initial policy or initial value function

